

TASO: Optimizing Deep Learning with Automatic Generation of Graph Substitutions

Zhihao Jia, Oded Padon, James Thomas, Todd Warszawski,
Matei Zaharia, and Alex Aiken

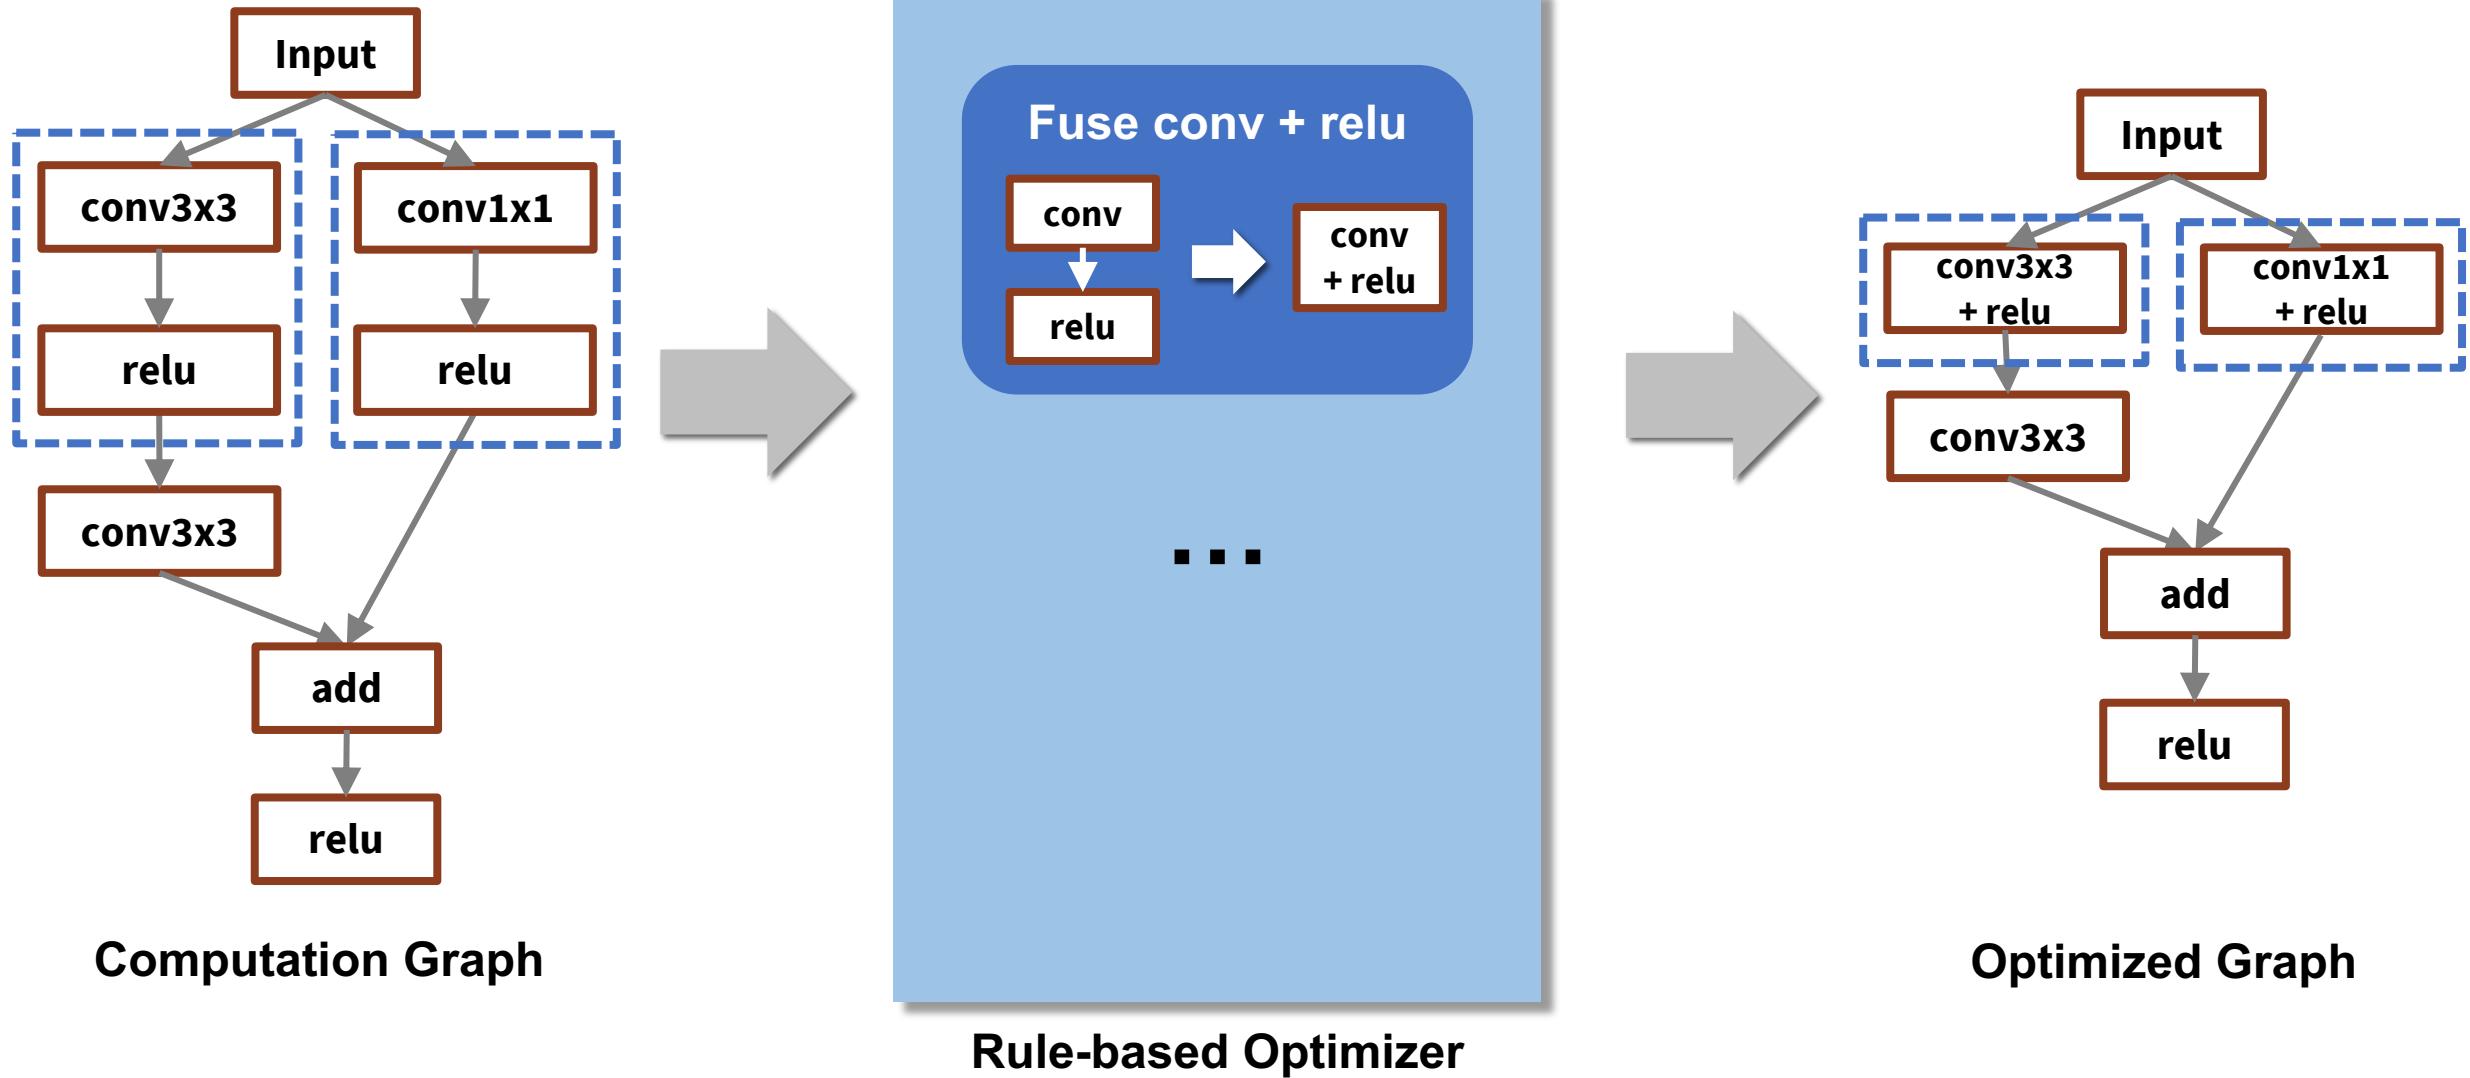
Stanford University

SOSP'19

12/14/19



Current Rule-based DNN Optimizations



Current Rule-based DNN Optimizations

TensorFlow currently includes ~200 rules (~53,000 LOC)

Fuse conv + relu

Fuse conv + batch normalization

Fuse multi. convs

...

Rule-based Optimizer

```
26 namespace tensorflow {
27 namespace graph_transforms {
28
29 // Converts Conv2D or MatMul ops followed by column-wise Muls into equivalent
30 // ops with the Mul baked into the convolution weights, to save computation
31 // during inference.
32 Status FoldBatchNorms(const GraphDef* input_graph_def,
33                      const TransformFuncContext& context,
34                      GraphDef* output_graph_def) {
35   GraphDef replaced_graph_def;
36   TF_RETURN_IF_ERROR(ReplaceMatchingOpTypes(
37     input_graph_def, // clang-format off
38     {"Mul"}, // mul_node
39     {
40       {"Conv2D|MatMul|DepthwiseConv2dNative", // conv_node
41        {"*"}, // input_node
42        {"Const"}, // weights_node
43      },
44      {"Const"}, // mul_values_node
45    },
46    {"Const"}, // mul_values_node
47  }, // clang-format on
48  [] (const NodeMatch& match, const std::set<string>& input_nodes,
49       const std::set<string>& output_nodes,
50       std::vector<NodeDef*>& new_nodes) {
51    // Find all the nodes we expect in the subgraph.
52    const NodeDef* mul_node = match.node;
53    const NodeDef* conv_node = match.inputs[0].node;
54    const NodeDef* input_node = match.inputs[0].inputs[0].node;
55    const NodeDef* weights_node = match.inputs[0].inputs[1].node;
56    const NodeDef* mul_values_node = match.inputs[1].node;
57
58    // Check that nodes that we use are not used somewhere else.
59    for (const auto& node : {conv_node, weights_node, mul_values_node}) {
60      if (output_nodes.count(node.name())) {
61        // Return original nodes.
62        new_nodes->insert(new_nodes->end(),
63                           {mul_node, conv_node, input_node, weights_node,
64                            mul_values_node});
65      }
66    }
67
68    Tensor weights = GetNodeTensorAttr(weights_node, "value");
69    Tensor mul_values = GetNodeTensorAttr(mul_values_node, "value");
70
71    // Make sure all the inputs really are vectors, with as many entries as
72    // there are columns in the weights.
73    int64 weights_cols;
74    if (conv_node.op() == "Conv2D") {
75      weights_cols = weights.shape().dim_size(3);
76    } else if (conv_node.op() == "DepthwiseConv2dNative") {
77      weights_cols =
78        weights.shape().dim_size(2) * weights.shape().dim_size(3);
79    } else {
80      weights_cols = weights.shape().dim_size(1);
81    }
82    if ((mul_values.shape().dims() != 1) ||
83        (mul_values.shape().dim_size(0) != weights_cols)) {
84      return errors::InvalidArgument(
85        "Mul constant input to batch norm has bad shape: ",
86        mul_values.DebugString());
87    }
88
89    // Multiply the original weights by the scale vector.
90    auto weights_vector = weights.flat<float>();
91    Tensor scaled_weights(DT_FLOAT, weights.shape());
92    auto scaled_weights_vector = scaled_weights.flat<float>();
93    for (int64 row = 0; row < weights_vector.dimension(0); ++row) {
94      scaled_weights_vector(row) =
95        weights_vector(row) *
96        mul_values.flat<float>()(row % weights_cols);
97    }
98
99    // Construct the new nodes.
100   NodeDef scaled_weights_node;
101   scaled_weights_node.set_op("Const");
102   scaled_weights_node.set_name(weights_node.name());
103   SetNodeAttr("dtype", DT_FLOAT, &scaled_weights_node);
104   SetNodeTensorAttr<float>("value", scaled_weights, &scaled_weights_node);
105   new_nodes->push_back(scaled_weights_node);
106
107   new_nodes->push_back(input_node);
108
109   NodeDef new_conv_node;
110   new_conv_node = conv_node;
111   new_conv_node.set_name(mul_node.name());
112   new_nodes->push_back(new_conv_node);
113
114   return Status::OK();
115 },
116 }, {},
117 {}, &replaced_graph_def);
118 *output_graph_def = replaced_graph_def;
119
120 return Status::OK();
121 }
122
123 REGISTER_GRAPH_TRANSFORM("fold_batch_norms", FoldBatchNorms);
124
125 } // namespace graph_transforms
126 } // namespace tensorflow
```

Limitations of Rule-based Optimizations

Robustness

Experts' heuristics do not apply to all DNNs/hardware

Horovod with XLA is slower than without XLA (Tensorflow 1.12) #713

Closed LiweiPeng opened this issue on Dec 19, 2018 · 2 comments

LiweiPeng commented on Dec 19, 2018

I have a distributed nmt model (Transformer-based, AdamOptimizer) using Horovod (0.15.1). When I turned on XLA under tensorflow 1.12, the training speed is about 20% slower instead of faster.

This result is sampled after training 1.5-hours and 4000 steps. I am using 4 V100 GPUs for the training.

Because my current software is tightly coupled with Horovod, I couldn't test whether this is Horovod related or not.

Does anyone have experience on whether this is expected?

tgaddair added the **question** label on Dec 19, 2018

New issue

Assignees: No one assigned

Labels: question

Milestone: No milestone

Notifications: Subscribe

You're not receiving notifications

When I turned on XLA (TensorFlow's graph optimizer), the training speed is **about 20% slower**.

Tensorflow XLA makes it slower?

I am writing a very simple tensorflow program with XLA enabled. Basically it's something like:

```
import tensorflow as tf
def ChainSoftMax(x, n)
    tensor = tf.nn.softmax(x)
    for i in range(n-1):
        tensor = tf.nn.softmax(tensor)
    return tensor
```

config = tf.ConfigProto()
config.graph_options.optimizer_options.global_jit_level = tf.OptimizerOptions.ON_1

```
input = tf.placeholder(tf.float32, [1000])
feed = np.random.rand(1000).astype('float32')

with tf.Session(config=config) as sess:
    res = sess.run(ChainSoftMax(input, 2000), feed_dict={input: feed})
```

Basically the idea is to see whether XLA can fuse the chain of softmax together to avoid multiple kernel launches. With XLA on, the above program is almost 2x slower than that without XLA on a machine with a GPU card. In my gpu profile, I saw XLA produces lots of kernels named as "reduce_xxx" and "fusion_xxx" which seem to overwhelm the overall runtime. Any one know what happened here?

With XLA, my program is **almost 2x slower than** without XLA

Limitations of Rule-based Optimizations

Robustness

Experts' heuristics do not apply to all DNNs/hardware

Scalability

New operators and graph structures require more rules

TensorFlow currently uses ~4K LOC to optimize convolution

Limitations of Rule-based Optimizations

Robustness

Experts' heuristics do not apply to all DNNs/hardware

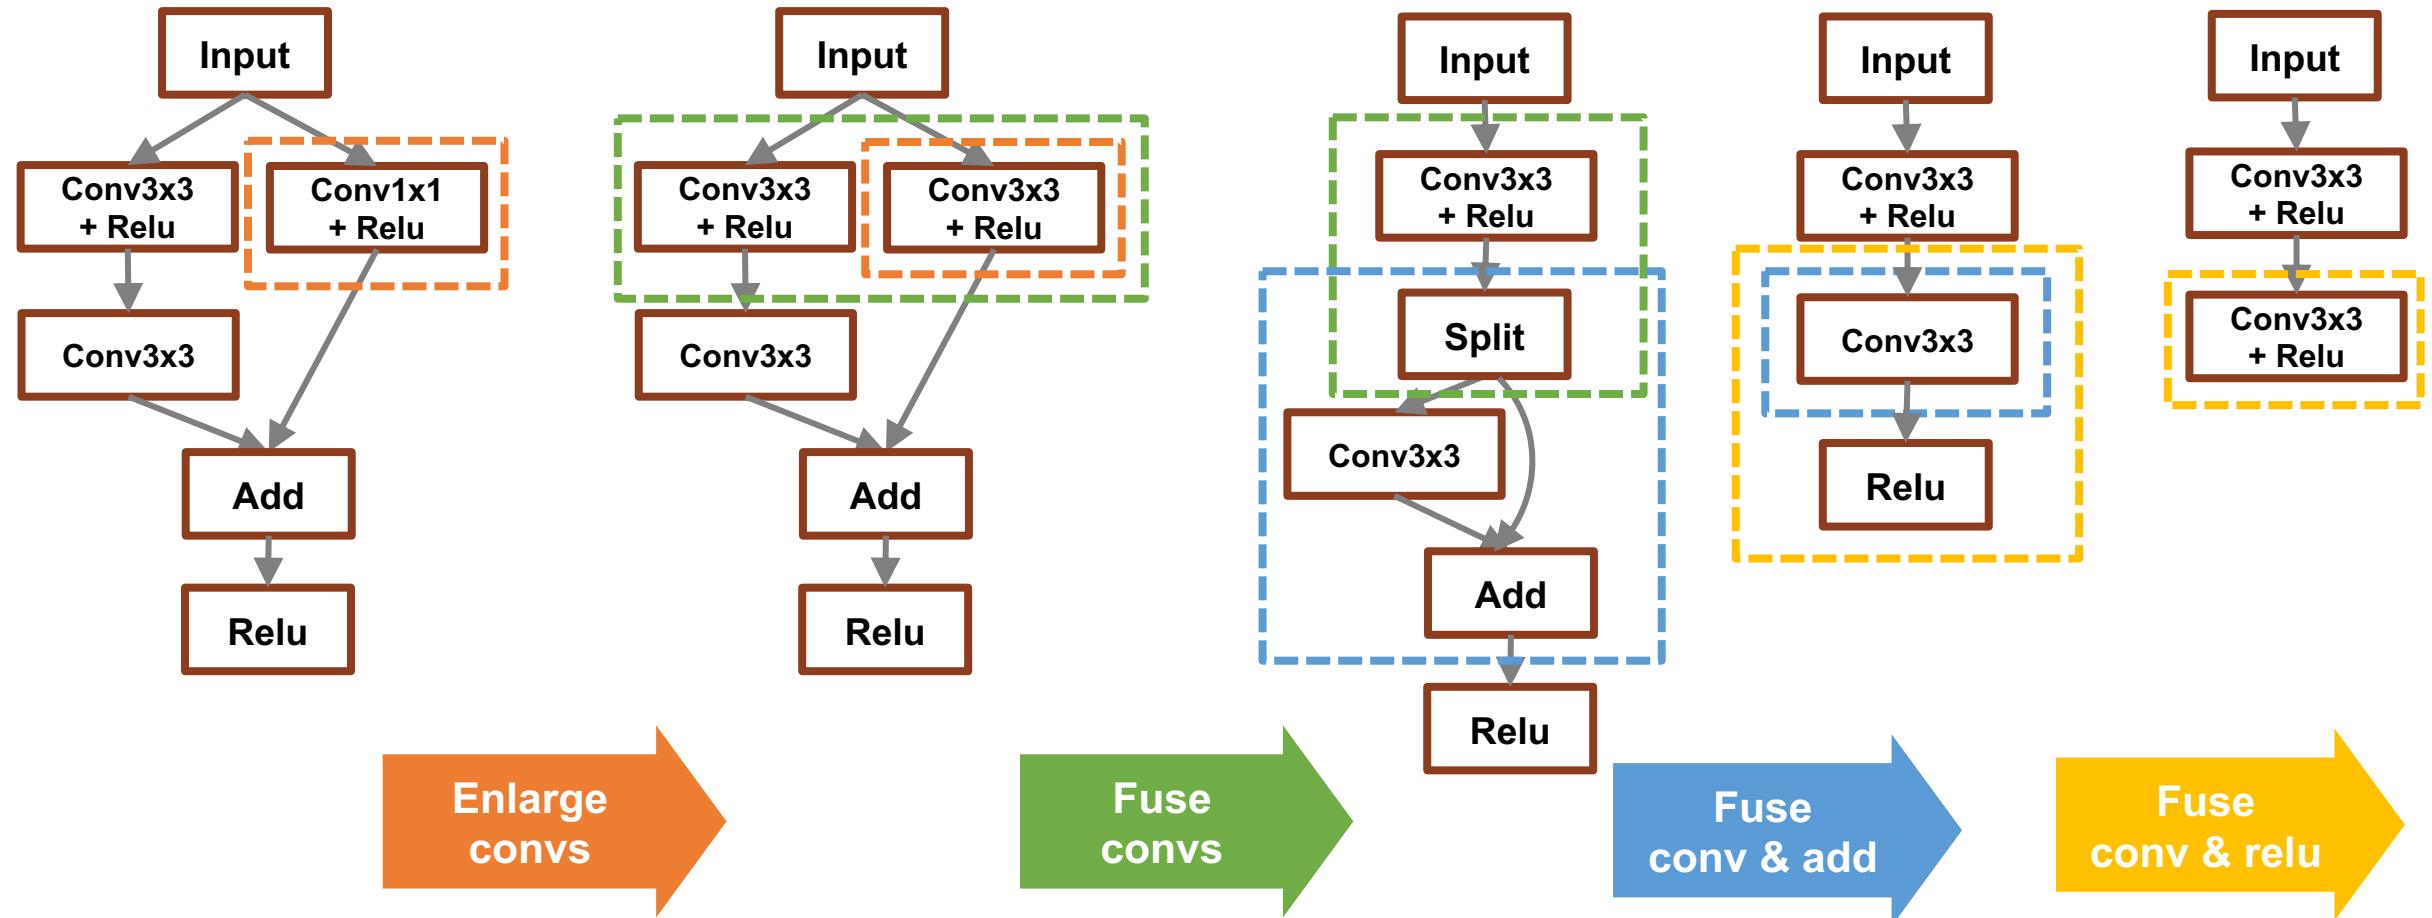
Scalability

New operators and graph structures require more rules

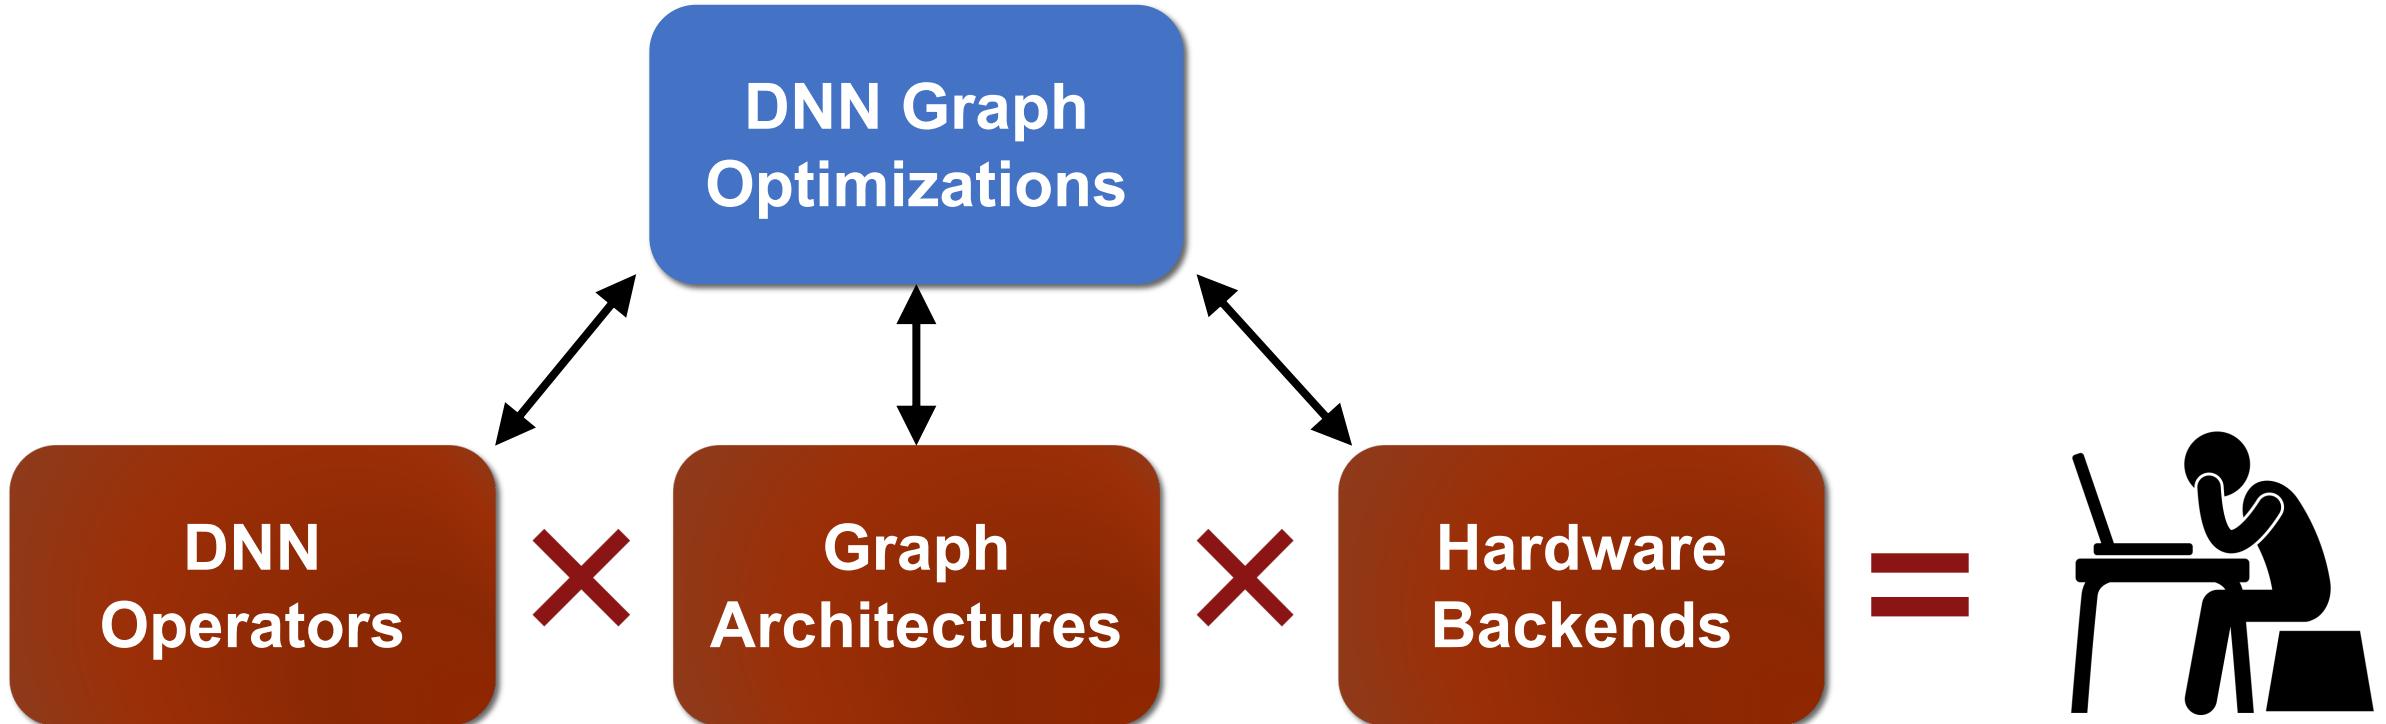
Performance

Miss subtle optimizations for specific DNNs/hardware

Motivating Example



The final graph is 30% faster on V100 but 10% slower on K80.

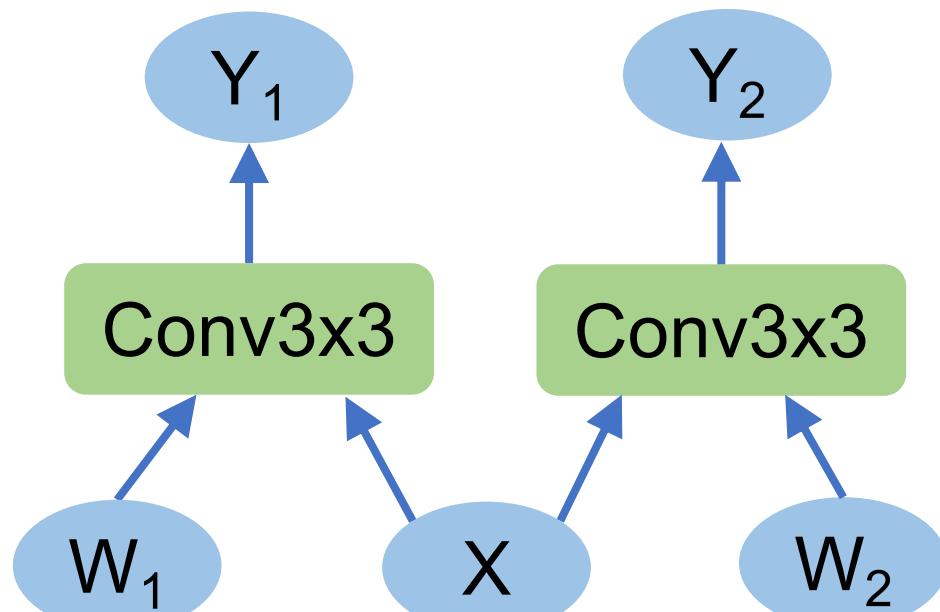


How should we address the complexity of designing DNN graph optimizations?

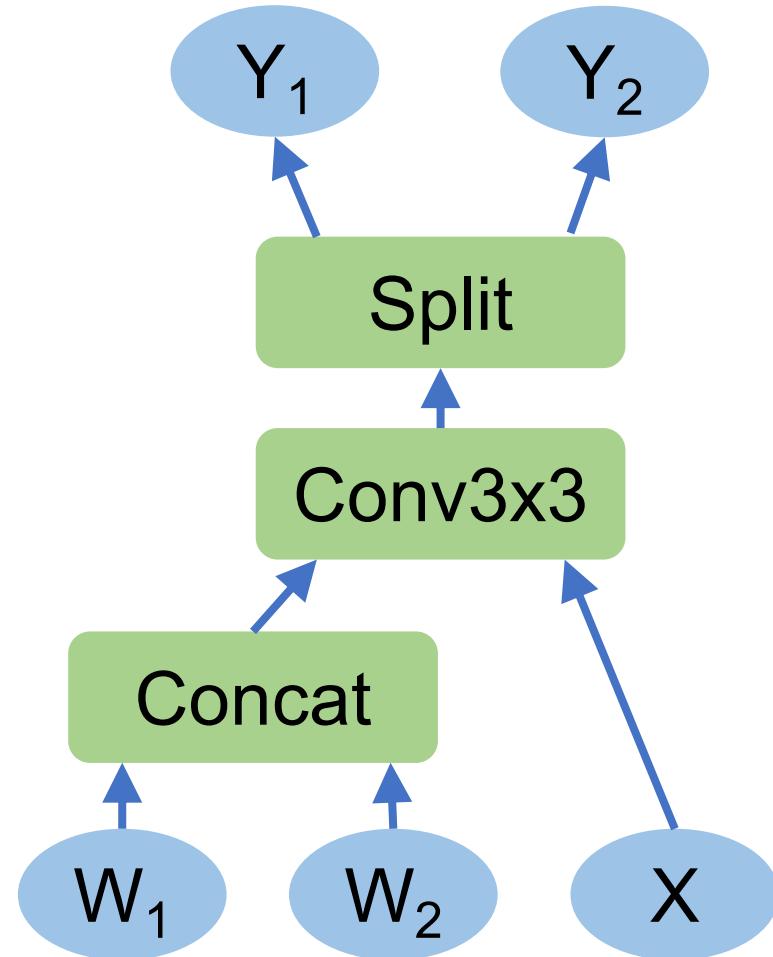
TASO: Tensor Algebra SuperOptimizer

- Key idea: replace manually-designed graph optimizations with ***automated generation and verification*** of graph substitutions for deep learning
- **Less engineering effort:** 53,000 LOC for manual graph optimizations in TensorFlow → 1,400 LOC in TASO
- **Better performance:** outperform existing optimizers by up to 2.8x

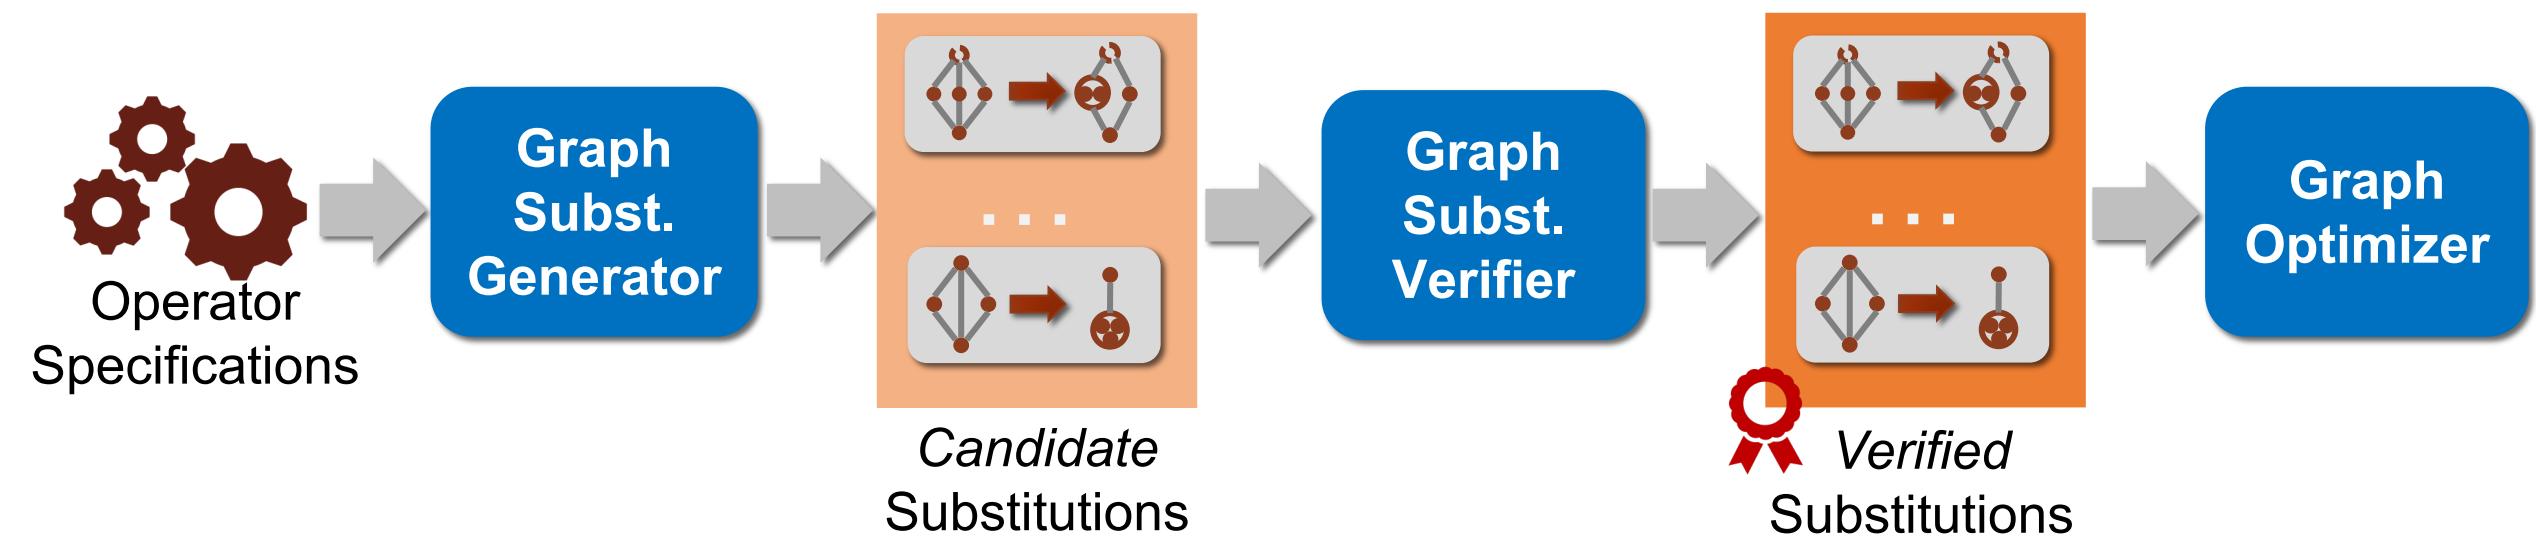
Graph Substitution



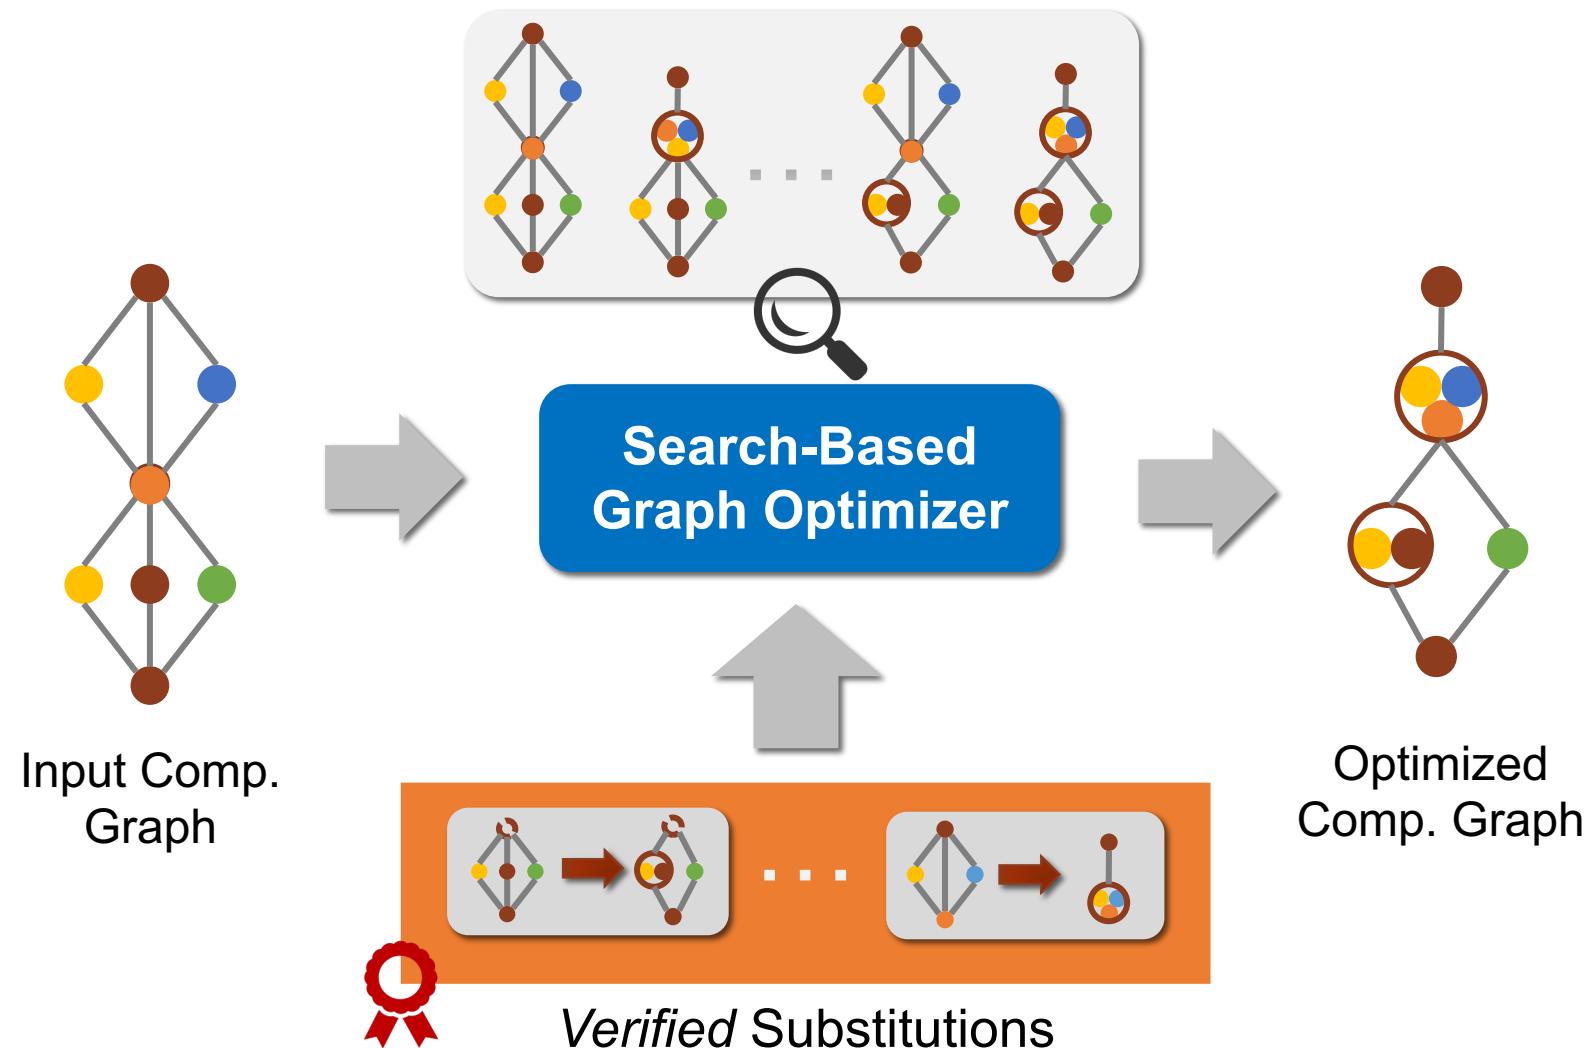
=



TASO Workflow



TASO Workflow



Key Challenges

1. How to generate potential substitutions?

Graph fingerprints

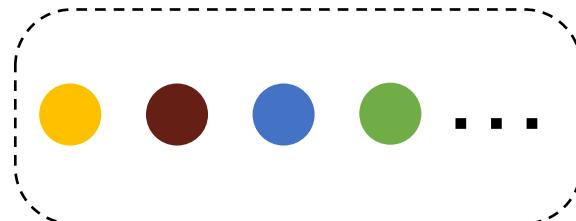
2. How to verify their correctness?

Operator specifications + theorem prover

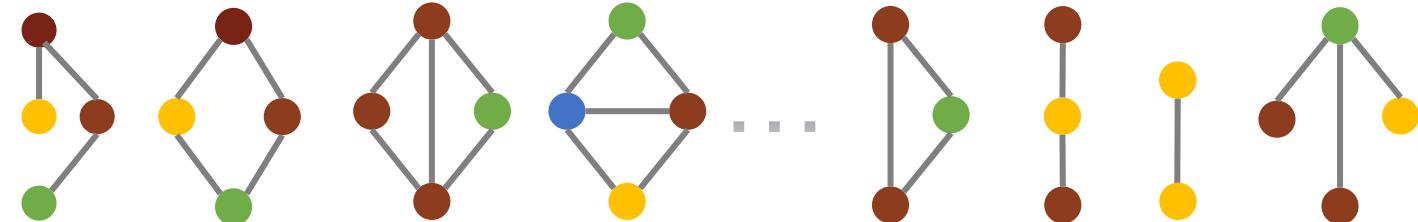
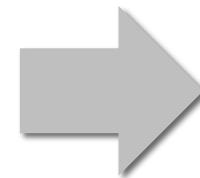
Graph Substitution Generator



Enumerate all possible graphs up to a fixed size using available operators



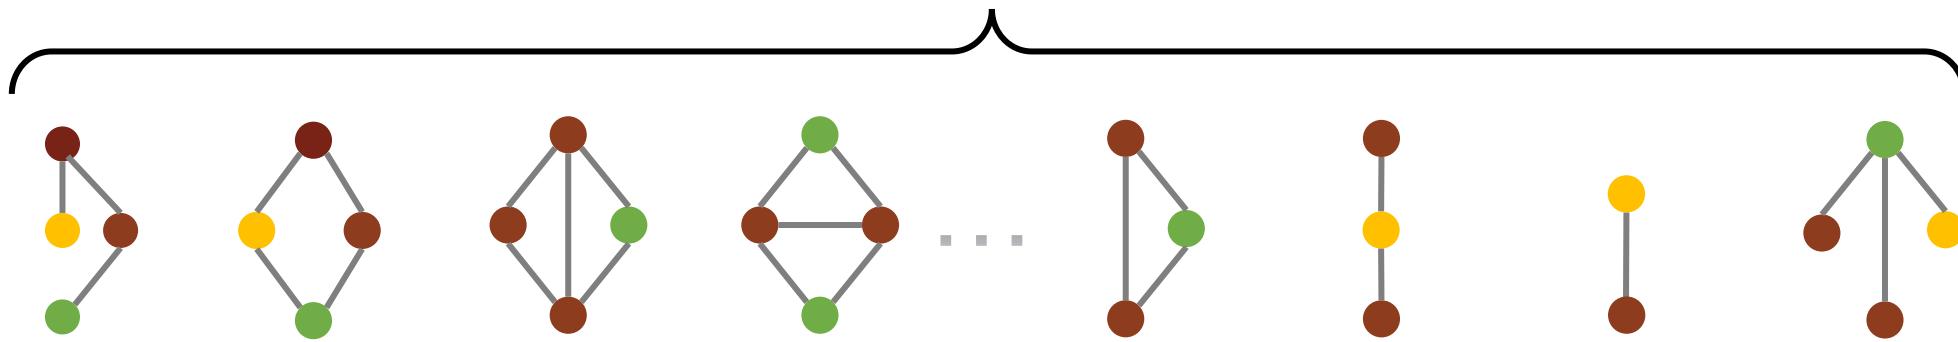
Operators supported by hardware backend



Graph Substitution Generator



66M graphs with up to 4 operators

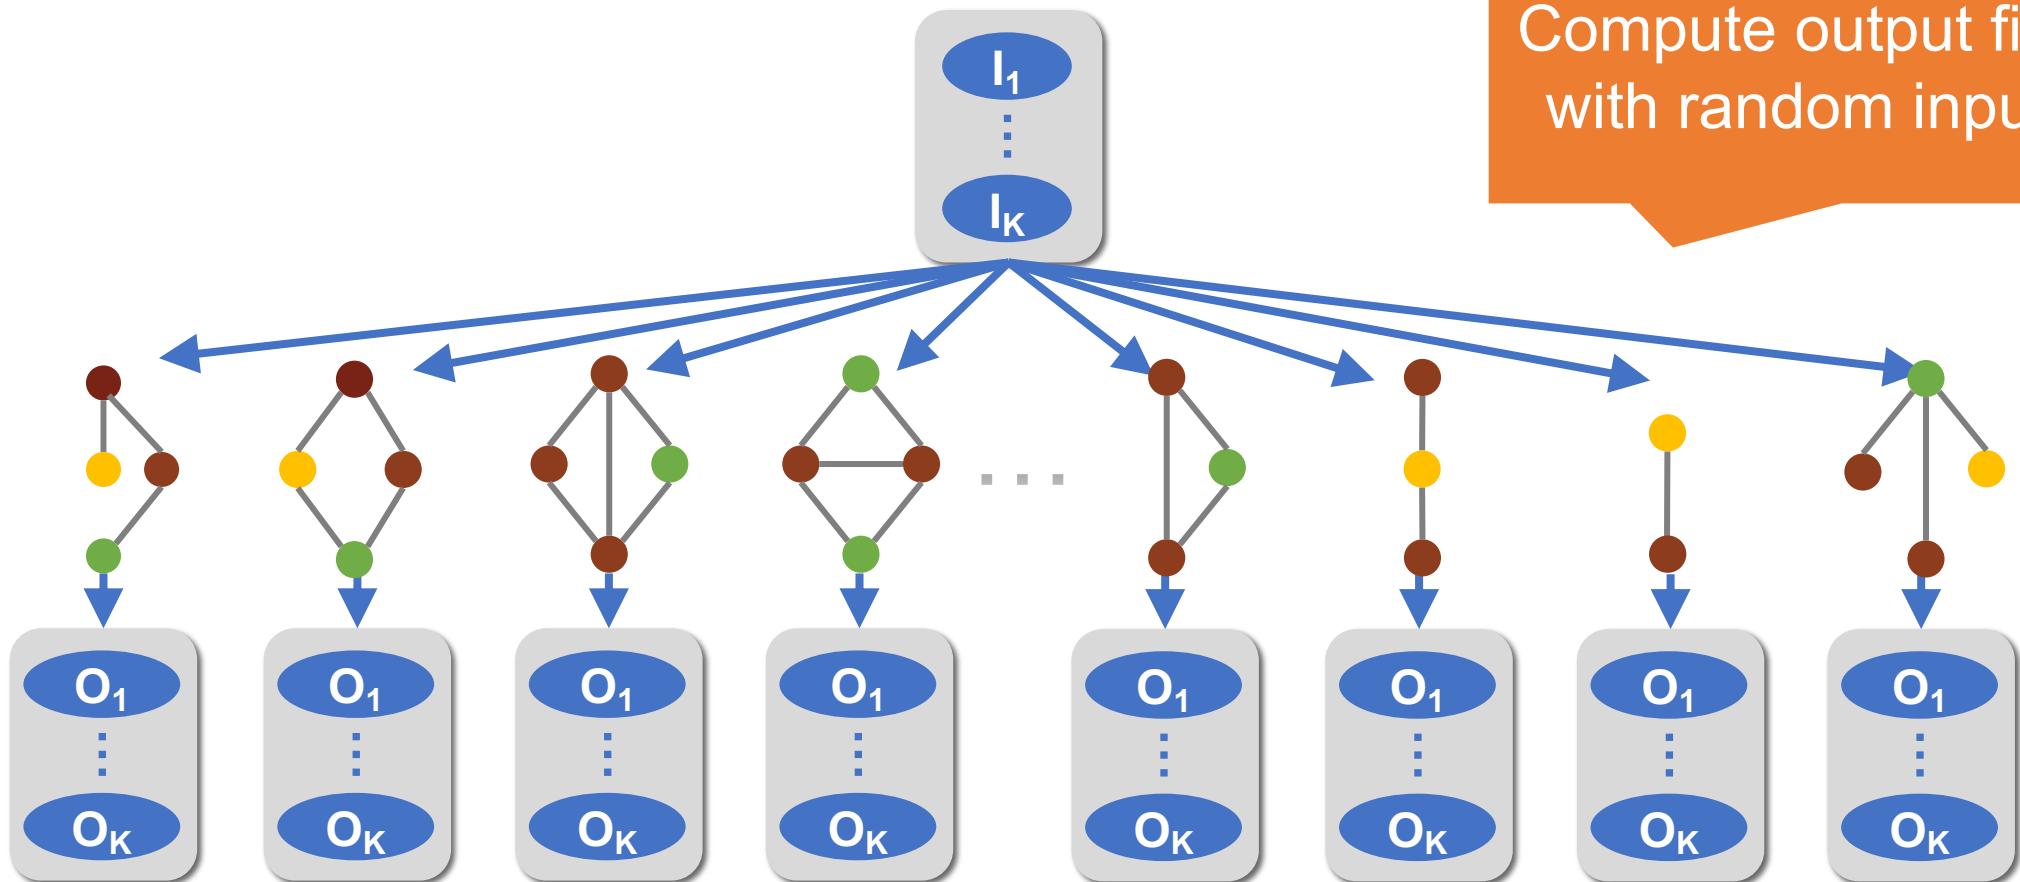


Directly evaluating all pairs requires a quadratic number of tests.

Graph Substitution Generator



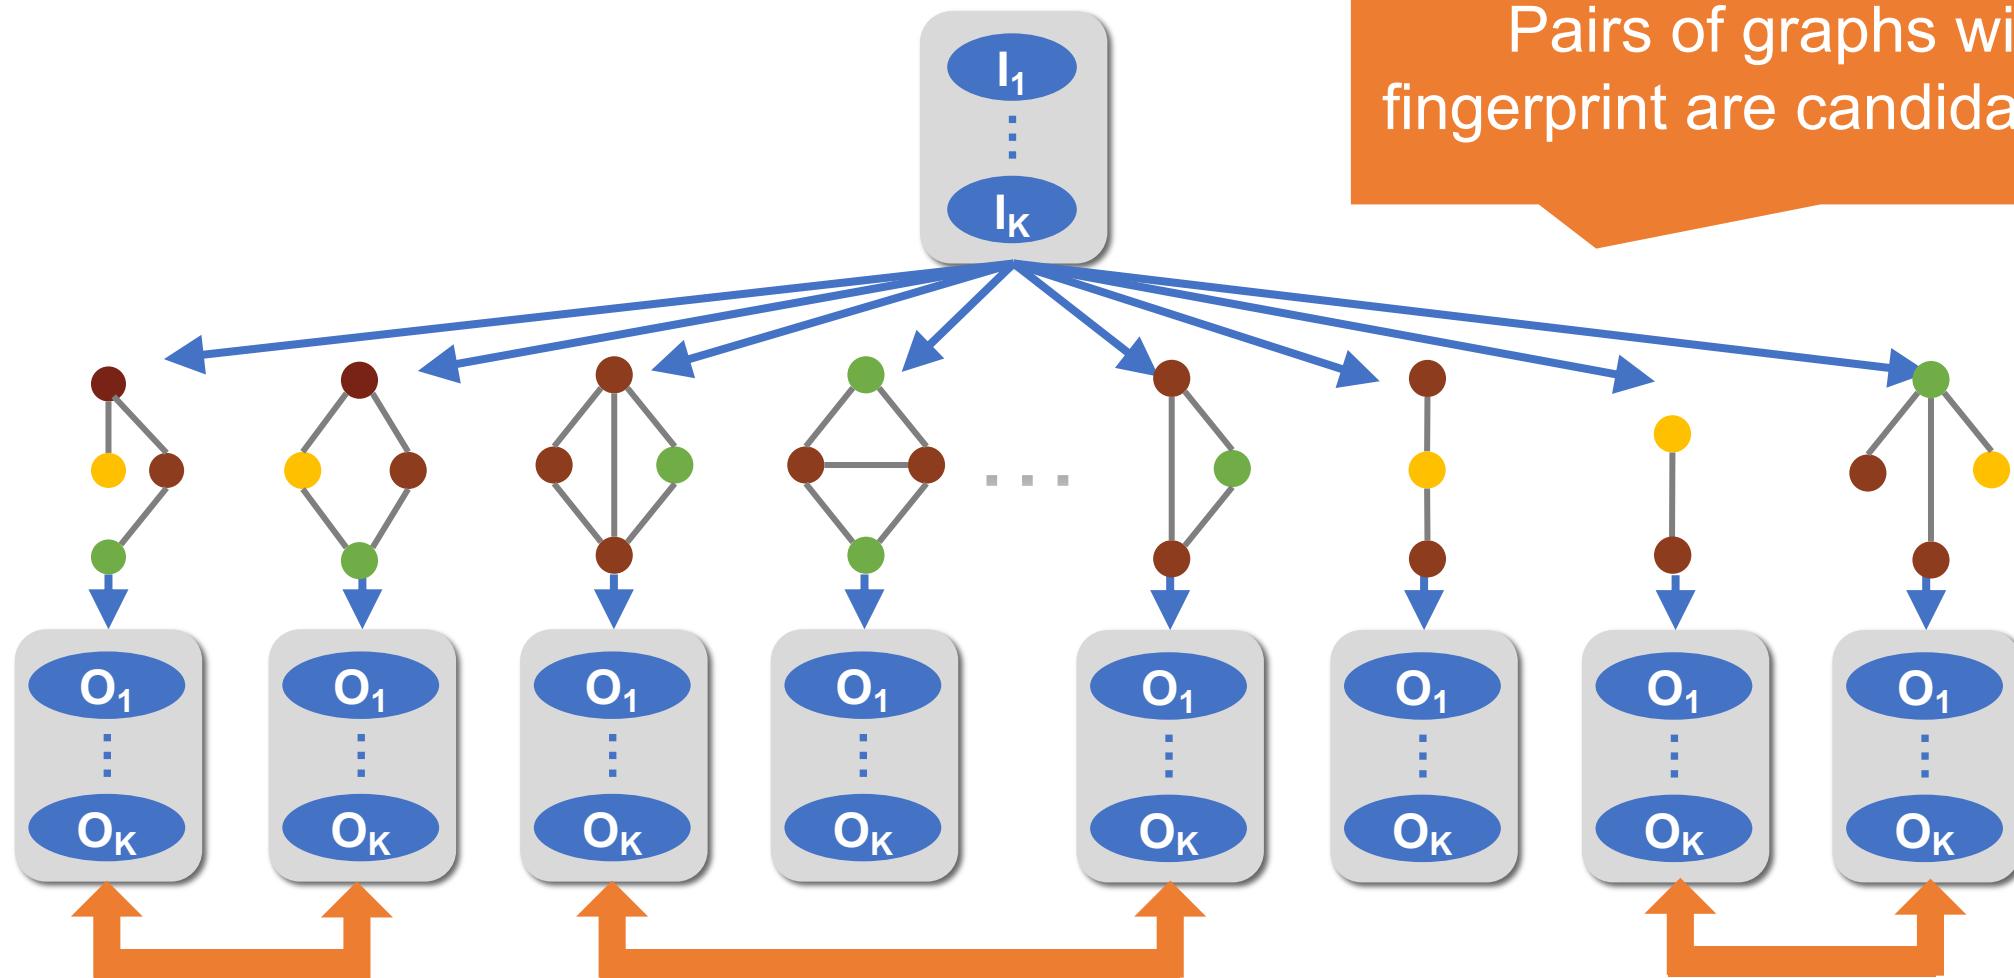
Compute output fingerprints
with random input tensors



Graph Substitution Generator



Pairs of graphs with identical
fingerprint are candidate substitutions



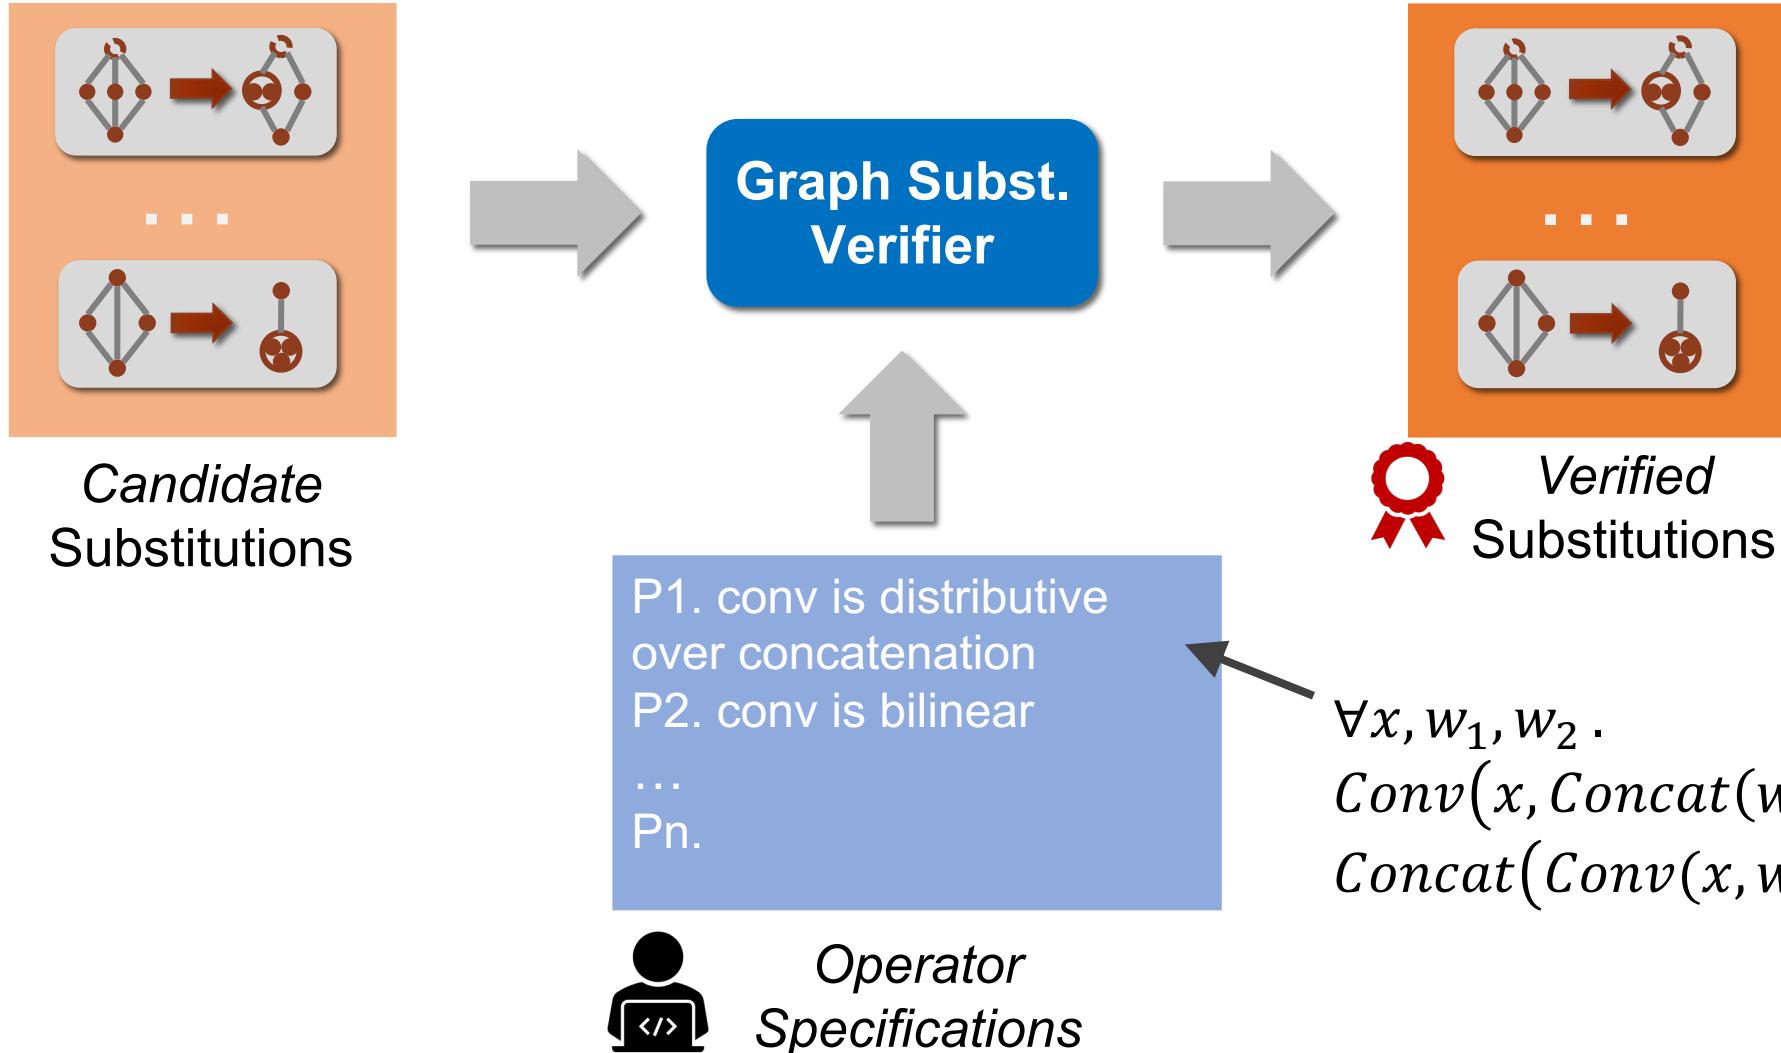
Graph Substitution Generator



TASO generates ~29,000 substitutions by enumerating graphs w/ up to 4 operators

743 substitutions remain after applying pruning techniques to eliminate redundancy

Graph Substitution Verifier



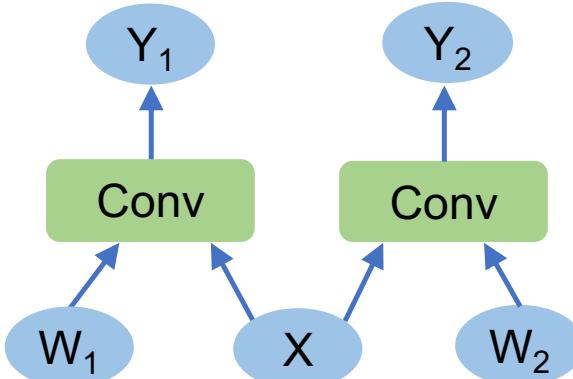
Verification Workflow

$\exists x, w_1, w_2 .$
 $(Conv(x, w_1), Conv(x, w_2))$
 $\neq Split(Conv(x, Concat(w_1, w_2)))$

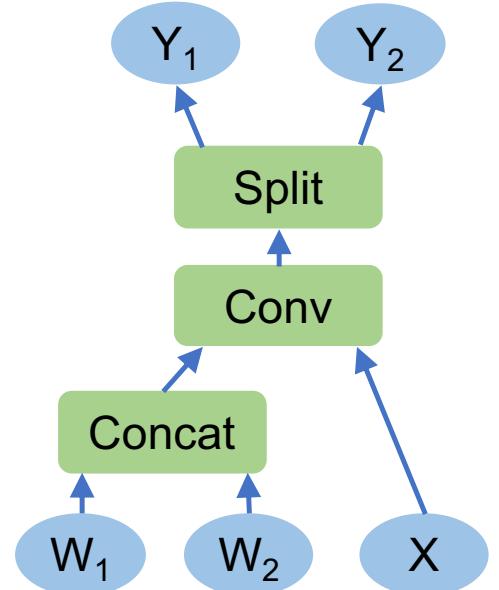
P1. $\forall x, w_1, w_2 .$
 $Conv(x, Concat(w_1, w_2)) =$
 $Concat(Conv(x, w_1), Conv(x, w_2))$

P2. ...

Operator Specifications



?
=



$(Conv(x, w_1), Conv(x, w_2))$

$Split(Conv(x, Concat(w_1, w_2)))$



Theorem
Prover

UNSAT

Verification Effort

TASO generates all 743 substitutions in 5 minutes, and verifies them against 43 operator properties in 10 minutes

Supporting a new operator requires a few hours of human effort to discover its properties

Operator specifications in TASO \approx 1,400 LOC
 Manual graph optimizations in TensorFlow \approx 53,000 LOC

| Operator Property | Comment |
|--|--------------------------------|
| $\forall x, y, z. \text{ewadd}(x, \text{ewadd}(y, z)) = \text{ewadd}(\text{ewadd}(x, y), z)$ | ewadd is associative |
| $\forall x, y. \text{ewadd}(x, y) = \text{ewadd}(y, x)$ | ewadd is commutative |
| $\forall x, y, z. \text{ewmul}(x, \text{ewmul}(y, z)) = \text{ewmul}(\text{ewmul}(x, y), z)$ | ewmul is associative |
| $\forall x, y. \text{ewmul}(x, y) = \text{ewmul}(y, x)$ | ewmul is commutative |
| $\forall x, y, z. \text{ewmul}(\text{ewadd}(x, y), z) = \text{ewadd}(\text{ewmul}(x, z), \text{ewmul}(y, z))$ | distributivity |
| $\forall x, y, w. \text{smul}(\text{smul}(x, y), w) = \text{smul}(x, \text{smul}(y, w))$ | smul is associative |
| $\forall x, y, w. \text{smul}(\text{ewadd}(x, y), w) = \text{ewadd}(\text{smul}(x, w), \text{smul}(y, w))$ | distributivity |
| $\forall s, p, x, y, w. \text{conv}_2(\text{conv}(s, p, \text{A}_{\text{none}}, x, y), w) = \text{conv}(s, p, \text{A}_{\text{none}}, \text{conv}_1(x, w), y)$ | conv is bilinear |
| $\forall s, p, x, y, z. \text{conv}(s, p, \text{A}_{\text{none}}, x, \text{ewadd}(y, z)) = \text{ewadd}(\text{conv}(s, p, \text{A}_{\text{none}}, x, y), \text{conv}(s, p, \text{A}_{\text{none}}, x, z))$ | conv is bilinear |
| $\forall s, p, c, x, y, z. \text{concat}_1(\text{conv}(s, p, c, x, y), \text{conv}(s, p, c, x, z)) = \text{conv}(s, p, c, x, \text{concat}(y, z))$ | convolution kernel |
| $\forall s, p, x, y. \text{relu}(\text{relu}(s, p, x, y)) = \text{relu}(s, p, \text{relu}(x, y))$ | A _{relu} applies relu |
| $\forall s, p, x, y, z. \text{conv}(s, p, \text{C}_{\text{pool}}, x, \text{ewadd}(y, z)) = \text{ewadd}(\text{conv}(s, p, \text{C}_{\text{pool}}, x, y), \text{conv}(s, p, \text{C}_{\text{pool}}, x, z))$ | conv. with C _{pool} |
| $\forall s, p, x, y, z. \text{conv}(s, p, \text{C}_{\text{pool}}, \text{concat}(x, y), z) = \text{concat}(\text{conv}(s, p, \text{C}_{\text{pool}}, x, z), \text{conv}(s, p, \text{C}_{\text{pool}}, y, z))$ | kernel |
| $\forall s, p, x, y, z. \text{conv}(s, p, \text{C}_{\text{pool}}, \text{mat}(x, y), z) = \text{mat}(\text{conv}(s, p, \text{C}_{\text{pool}}, x, z), \text{conv}(s, p, \text{C}_{\text{pool}}, y, z))$ | matrix |
| $\forall a, x, y. \text{split}_0(a, \text{concat}(a, x, y)) = x$ | identity |
| $\forall a, x, y. \text{split}_1(a, \text{concat}(a, x, y)) = y$ | split definition |
| $\forall a, x, y. \text{concat}_0(a, \text{split}_0(a, x)) = x$ | definition |
| $\forall a, x, y. \text{concat}_1(a, \text{split}_1(a, x)) = y$ | of concatenation |
| $\forall a, x, y, z. \text{concat}_0(a, \text{concat}(x, y)) = \text{concat}(x, \text{concat}(a, y))$ | commutativity |
| $\forall a, x, y, z. \text{concat}_1(a, \text{concat}(x, y)) = \text{concat}(\text{concat}(a, x), y)$ | commutativity |
| $\forall a, x, y, z. \text{concat}_0(a, \text{concat}_0(x, y)) = \text{concat}_0(\text{concat}_0(a, x), y)$ | commutativity |
| $\forall a, x, y, z. \text{concat}_1(a, \text{concat}_1(x, y)) = \text{concat}_1(\text{concat}_1(a, x), y)$ | commutativity |
| $\forall a, x, y, z. \text{concat}_0(a, \text{concat}_0(x, y), z) = \text{concat}_0(\text{concat}_0(a, x), \text{concat}_0(y, z))$ | concatenation and transpose |
| $\forall a, x, y, z. \text{concat}_1(a, \text{concat}_1(x, y), z) = \text{concat}_1(\text{concat}_1(a, x), \text{concat}_1(y, z))$ | concatenation and matrix mul. |
| $\forall a, x, y, z. \text{concat}_0(a, \text{concat}_1(x, y), z) = \text{concat}_1(\text{concat}_0(a, x), \text{concat}_1(y, z))$ | concatenation and matrix mul. |
| $\forall a, x, y, z. \text{concat}_1(a, \text{concat}_0(x, y), z) = \text{concat}_0(\text{concat}_1(a, x), \text{concat}_0(y, z))$ | concatenation and conv. |
| $\forall s, p, x, y, z. \text{concat}_0(\text{conv}(s, p, \text{C}_{\text{conv}}, x, y), z) = \text{conv}(s, p, \text{C}_{\text{conv}}, \text{concat}(x, y), z)$ | concatenation and conv. |
| $\forall s, p, x, y, z. \text{concat}_1(\text{conv}(s, p, \text{C}_{\text{conv}}, x, y), z) = \text{conv}(s, p, \text{C}_{\text{conv}}, \text{concat}(1, \text{concat}(x, y)), z)$ | concatenation and conv. |
| $\forall k, s, p, x, y. \text{concat}(1, \text{pool}_{\text{avg}}(k, s, p, x), \text{pool}_{\text{avg}}(k, s, p, y)) = \text{pool}_{\text{avg}}(k, s, p, \text{concat}(1, \text{concat}(x, y)))$ | concatenation and pooling |
| $\forall k, s, p, x, y. \text{concat}(0, \text{pool}_{\text{max}}(k, s, p, x), \text{pool}_{\text{max}}(k, s, p, y)) = \text{pool}_{\text{max}}(k, s, p, \text{concat}(0, \text{concat}(x, y)))$ | concatenation and pooling |
| $\forall k, s, p, x, y. \text{concat}(1, \text{pool}_{\text{max}}(k, s, p, x), \text{pool}_{\text{max}}(k, s, p, y)) = \text{pool}_{\text{max}}(k, s, p, \text{concat}(1, \text{concat}(x, y)))$ | concatenation and pooling |

Search-Based Graph Optimizer¹

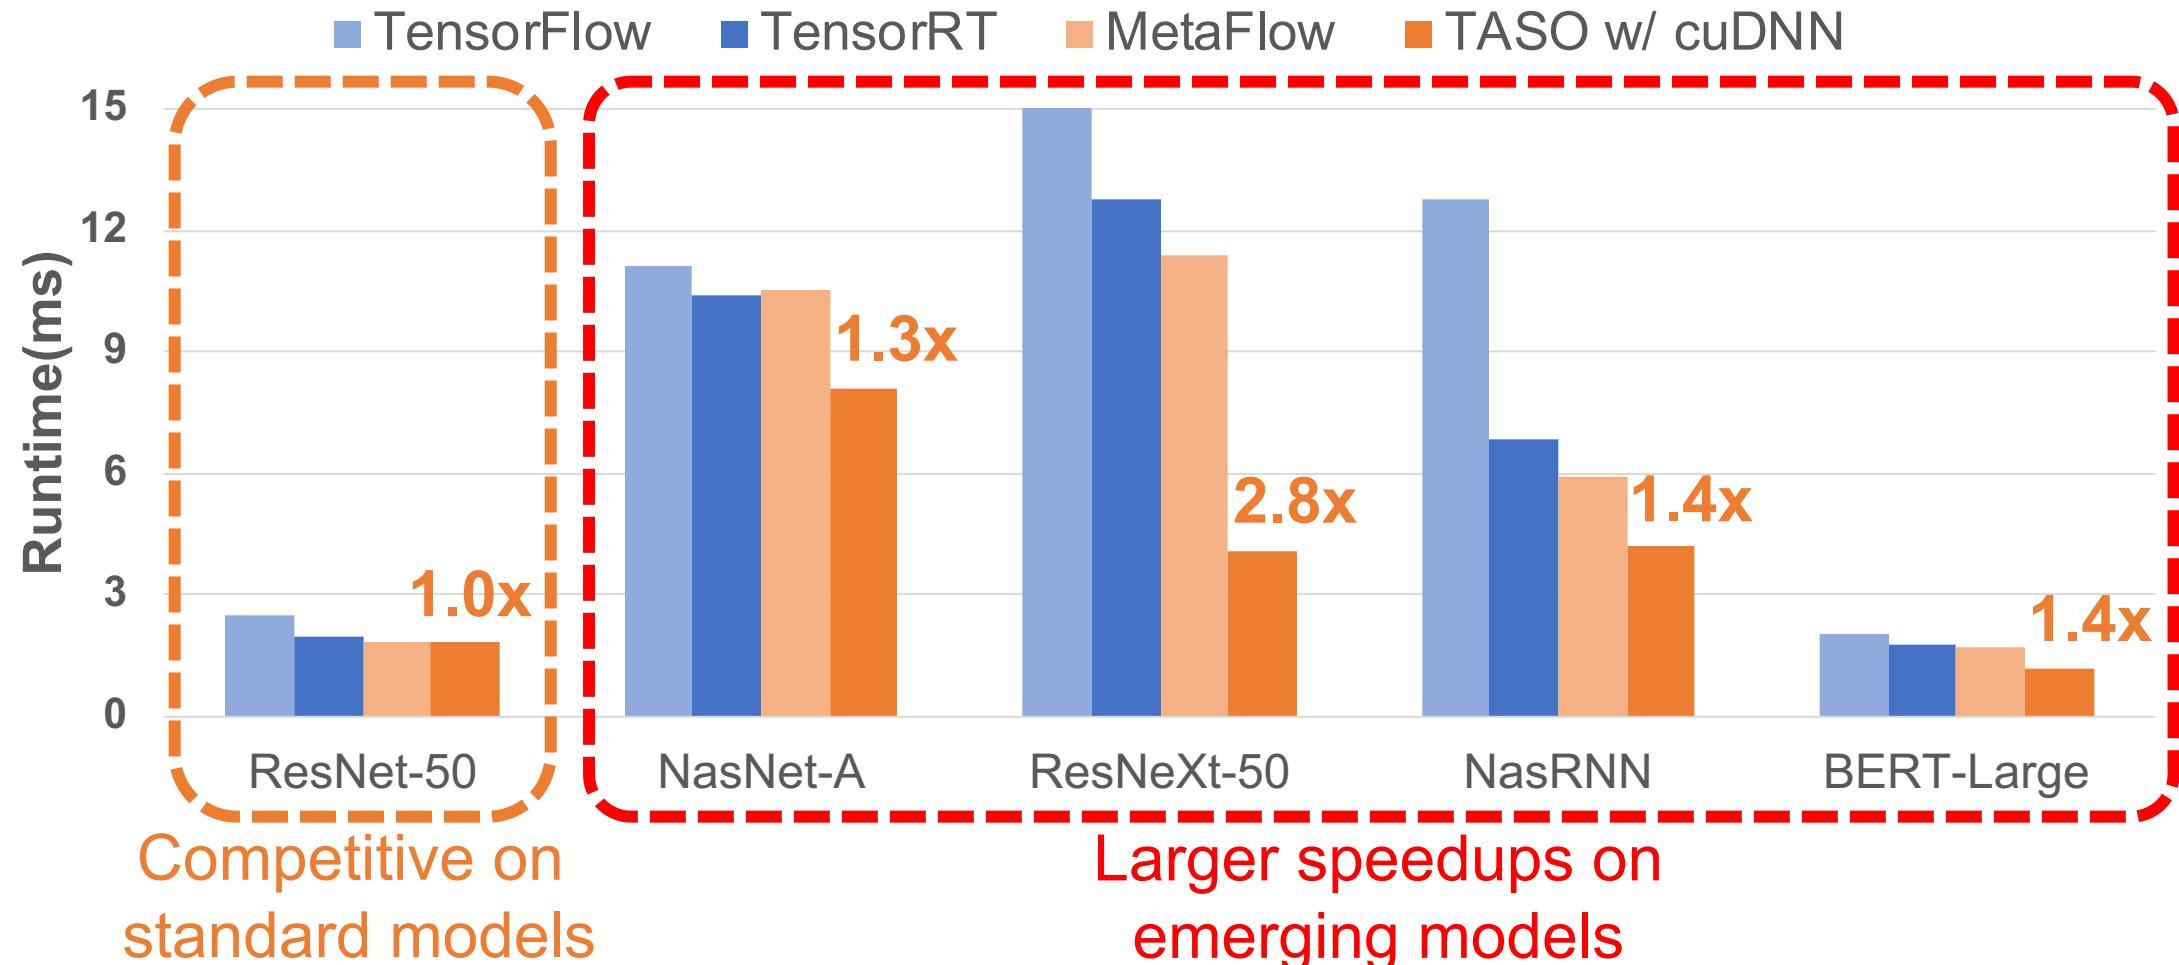


- **Goal:** applying verified substitutions to obtain an optimized graph
- **Cost model²**
 - Based on the sum of individual operators' cost
 - Measure the cost of each operator on hardware
- **Cost-based backtracking search**
 - Backtrack local optimal solutions
 - Optimizing a DNN model takes less than 10 minutes

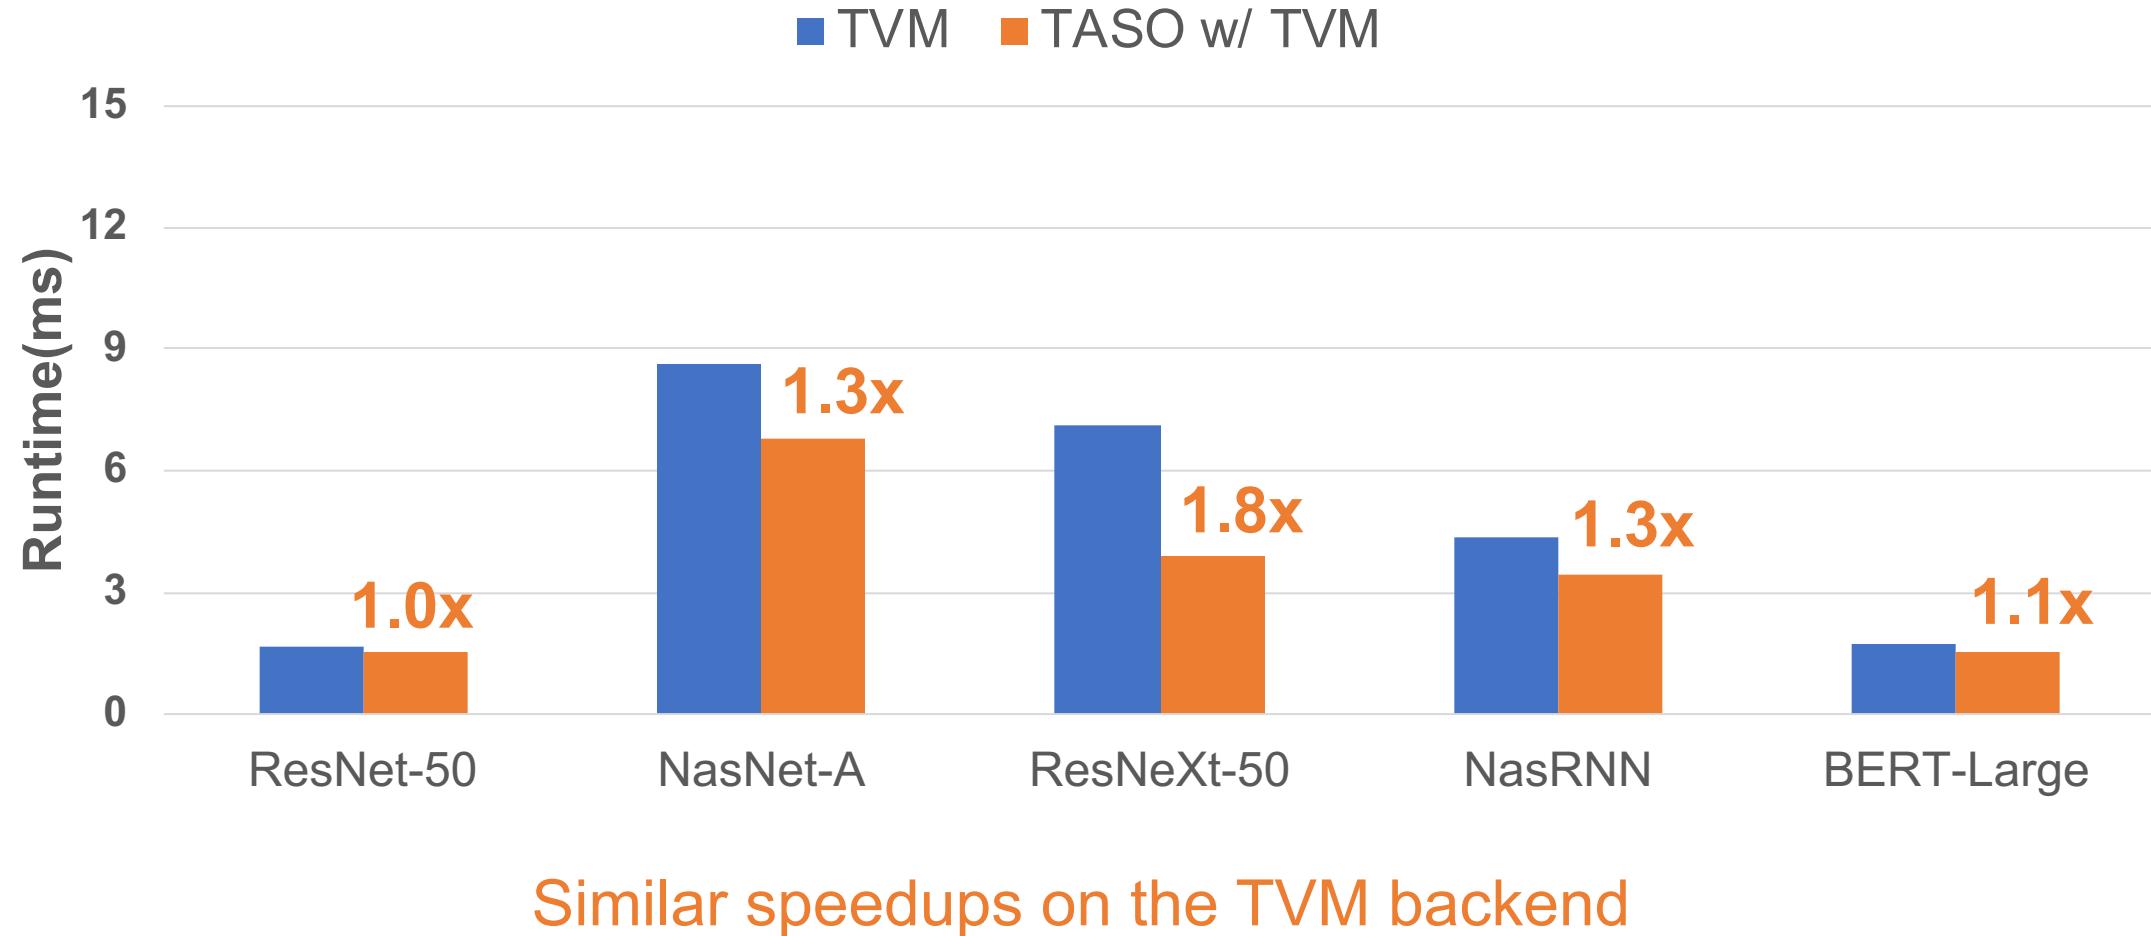
1. Z. Jia et al. Optimizing DNN Computation with Relaxed Graph Substitutions. In SysML'19.

2. Z. Jia et al. Exploring Hidden Dimensions in Parallelizing Convolutional Neural Networks. ICML'18.

End-to-end Inference Performance (V100 GPU w/ cuDNN)



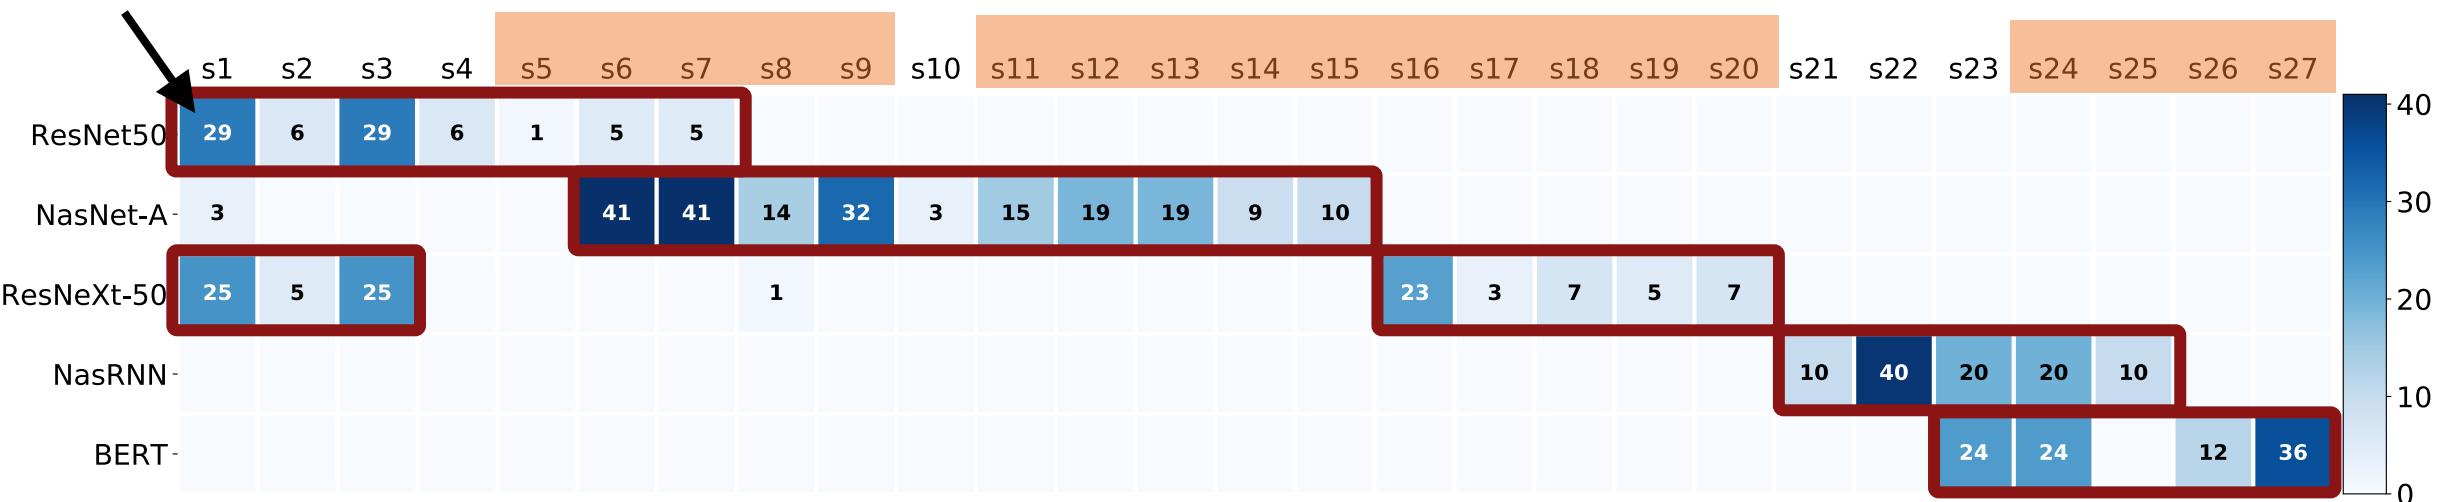
End-to-end Inference Performance (V100 GPU w/ TVM)



Heatmap of Used Substitutions

How many times a subst. is used to optimize a DNN

Not covered in TensorFlow



Different DNN models require **different substitutions**.

Conclusion

TASO is the first DNN optimizer that automatically generates substitutions

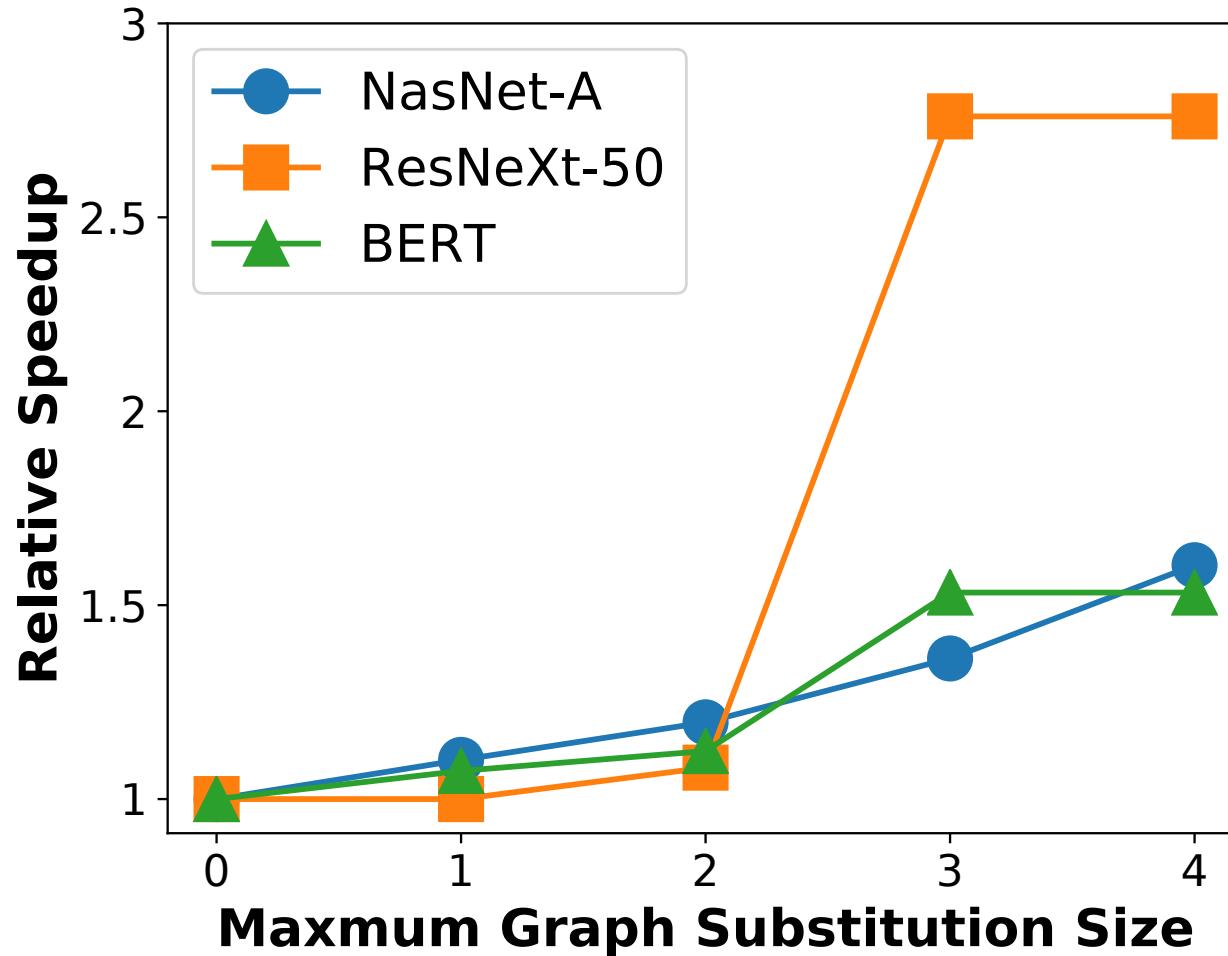
- Less engineering effort
- Better performance
- Formal verification

<https://github.com/jiazhihao/taso>

- Support DNN models in ONNX, TensorFlow, and PyTorch



Scalability Analysis

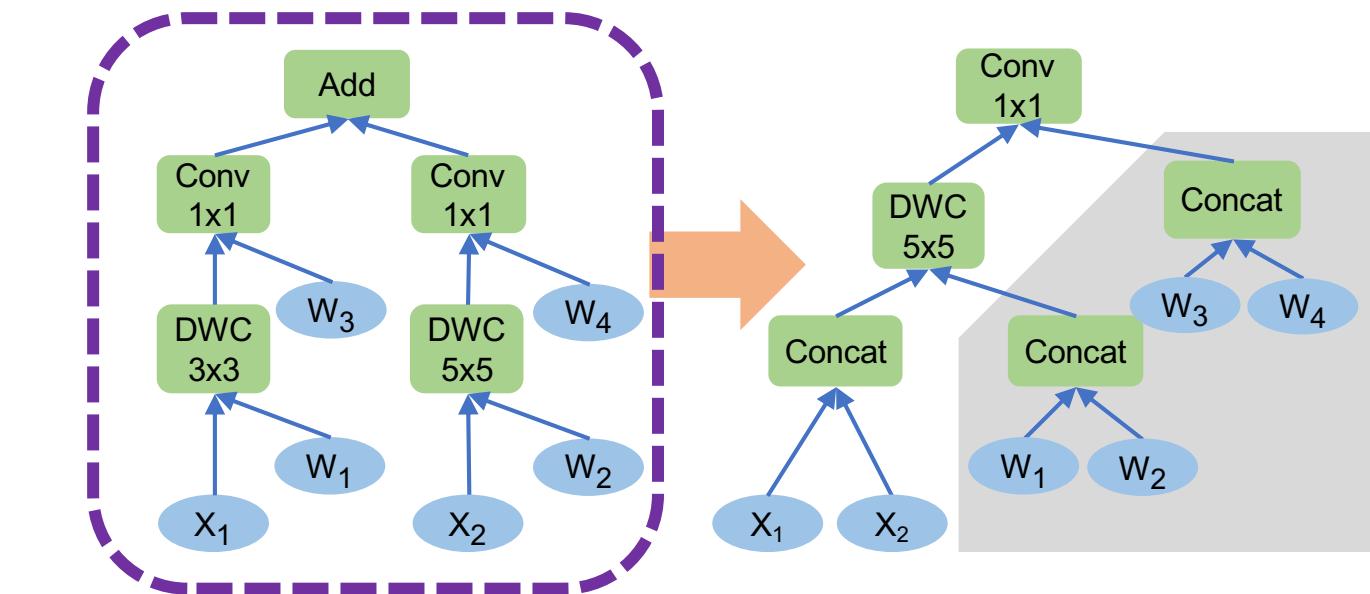
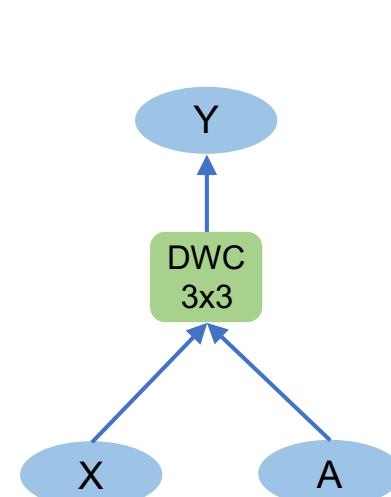
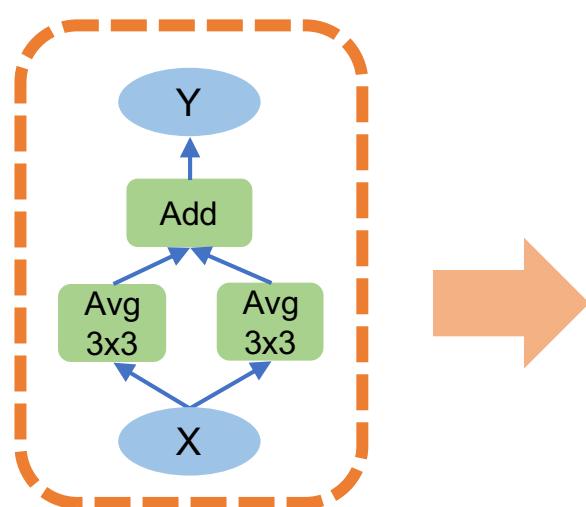
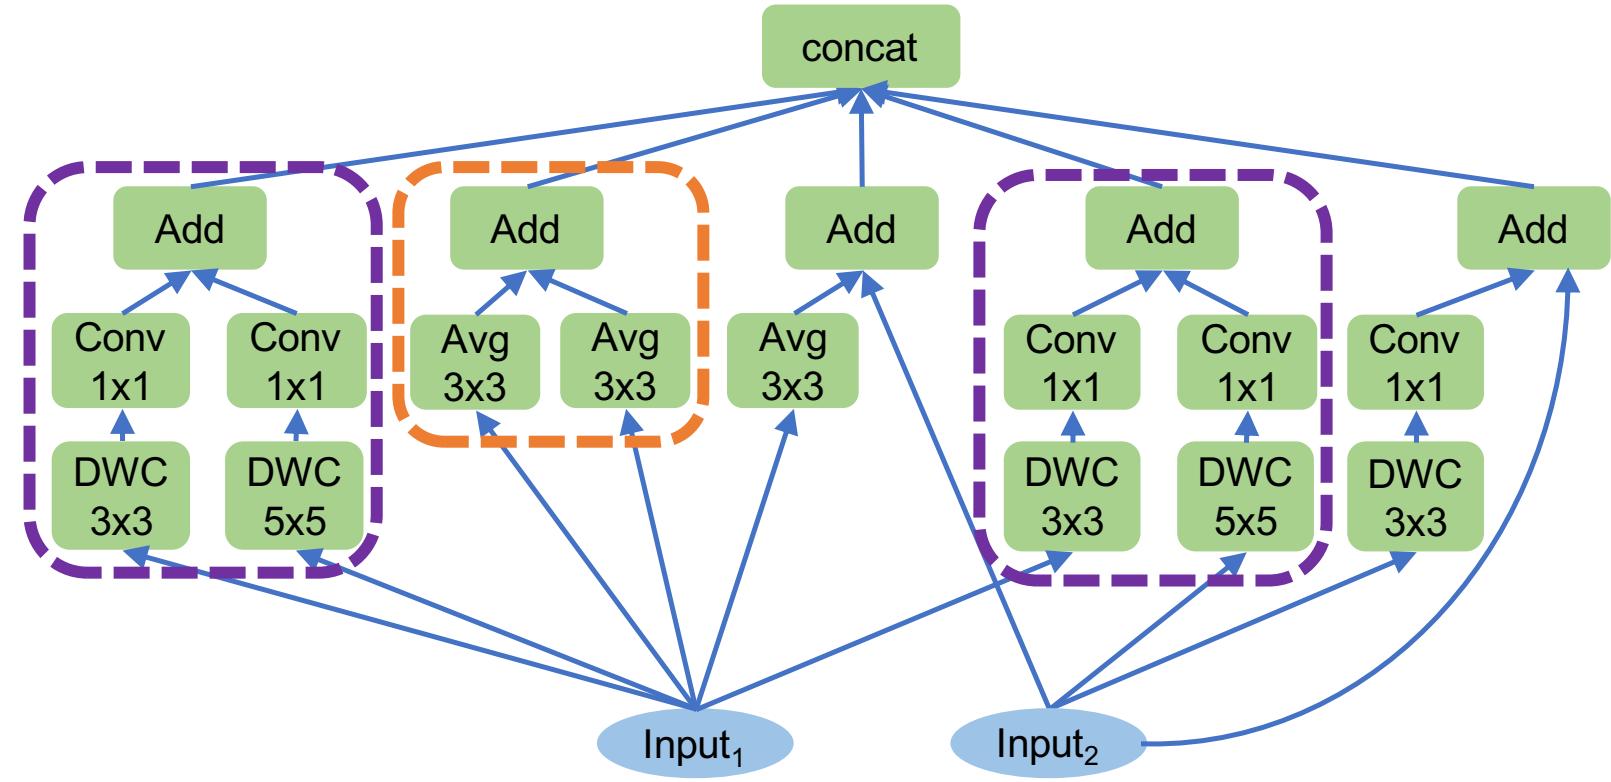


Case Study: NASNet

Add: element-wise addition

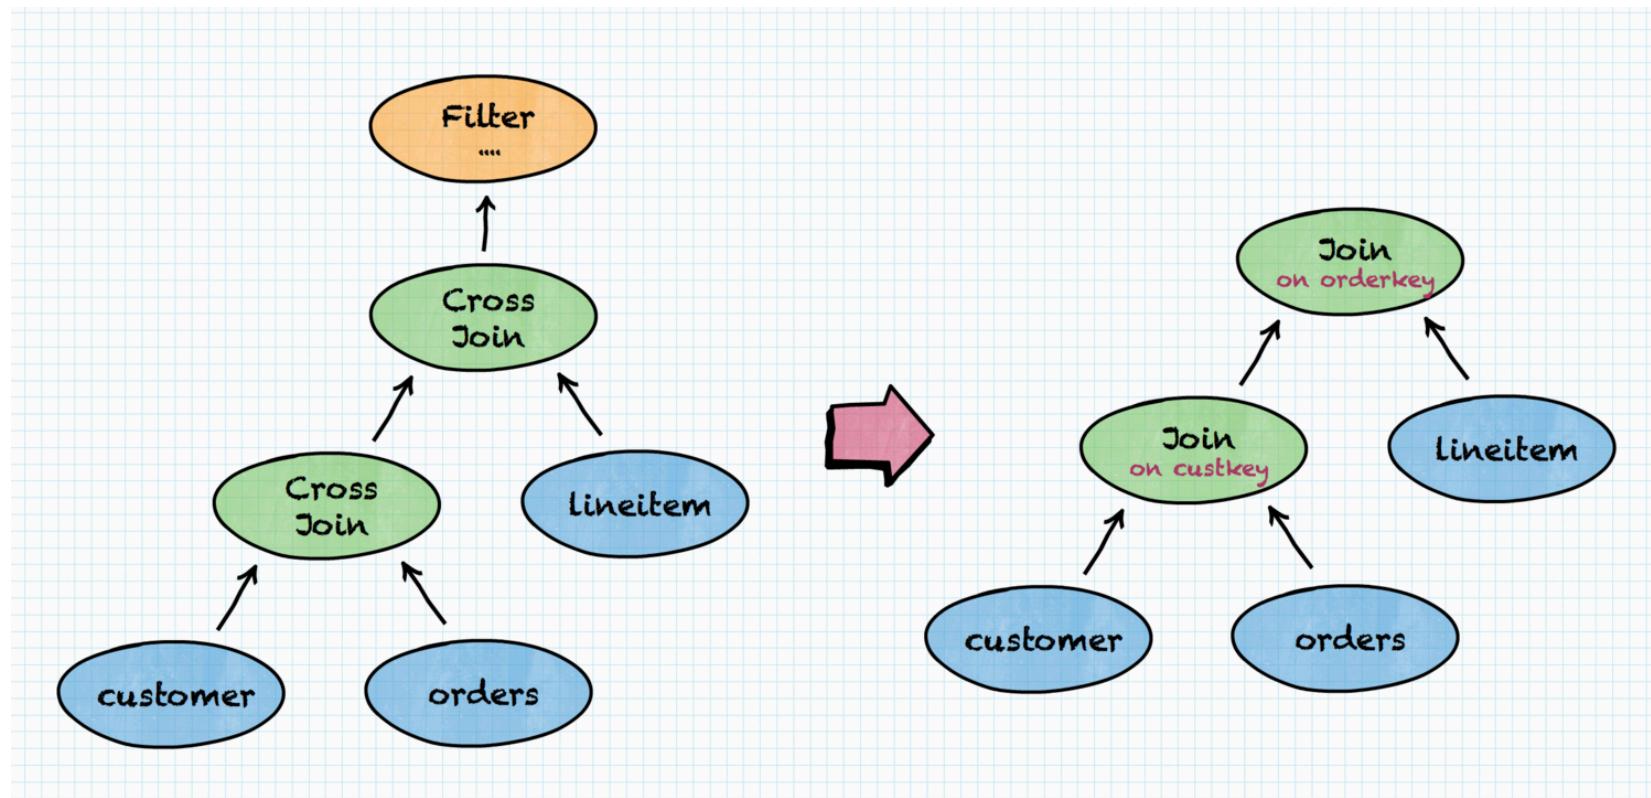
Conv: standard conv

DWC: depth-wise conv



Future Work: Query Optimizations

- A database query is expressed as a tree of relational operators
- Query optimizations are tree transformations



Contribution

- Replacing current manually-designed graph optimizations with ***automatic generation*** of graph substitutions for deep learning
- **Less engineering effort:** 53,000 LOC for graph optimizations in TensorFlow → 1,400 LOC
- **Better performance:** outperform existing optimizers by up to 2.8x
- **Correctness:** formal verification of graph substitutions

Limitations of Rule-based Optimizations

Robustness

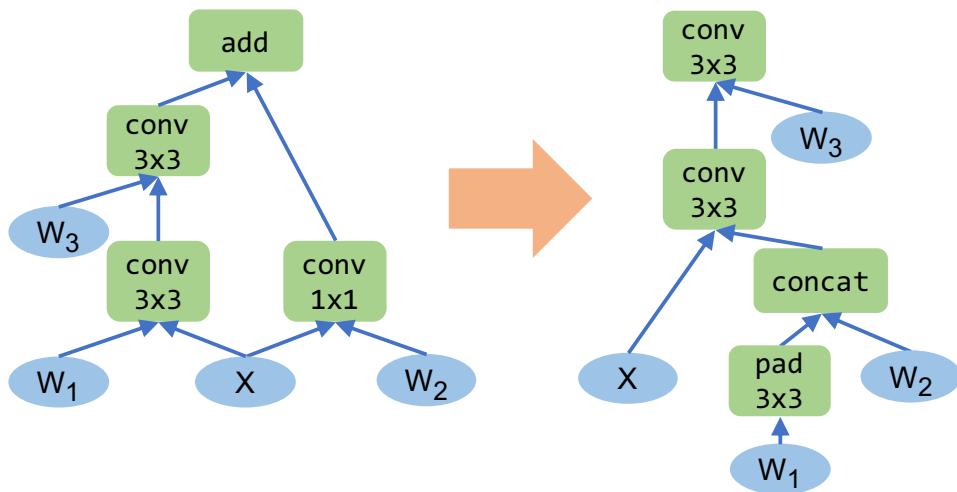
Experts' heuristics do not apply to all DNNs/hardware

Scalability

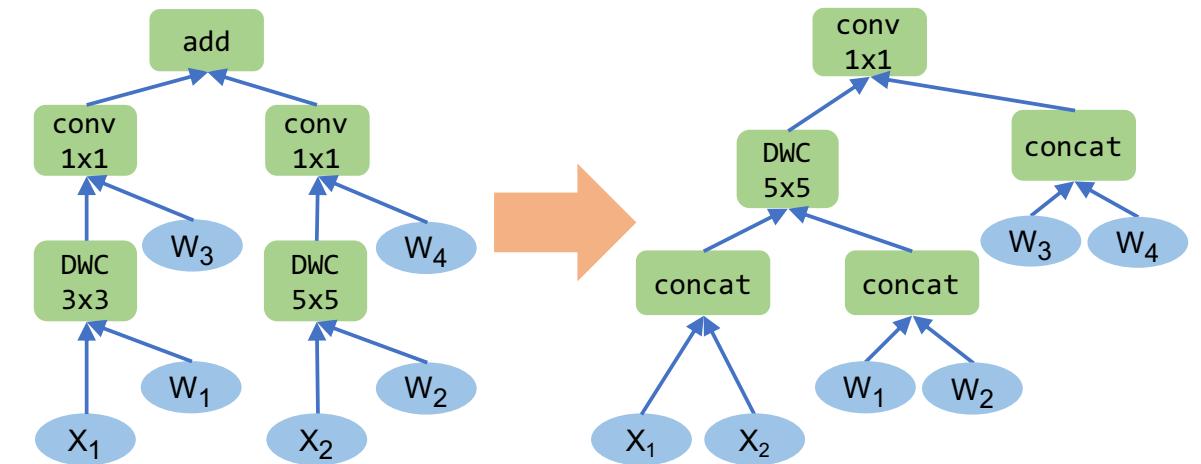
New operators and graph structures require more rules

Performance

Miss subtle optimizations for specific DNNs/hardware



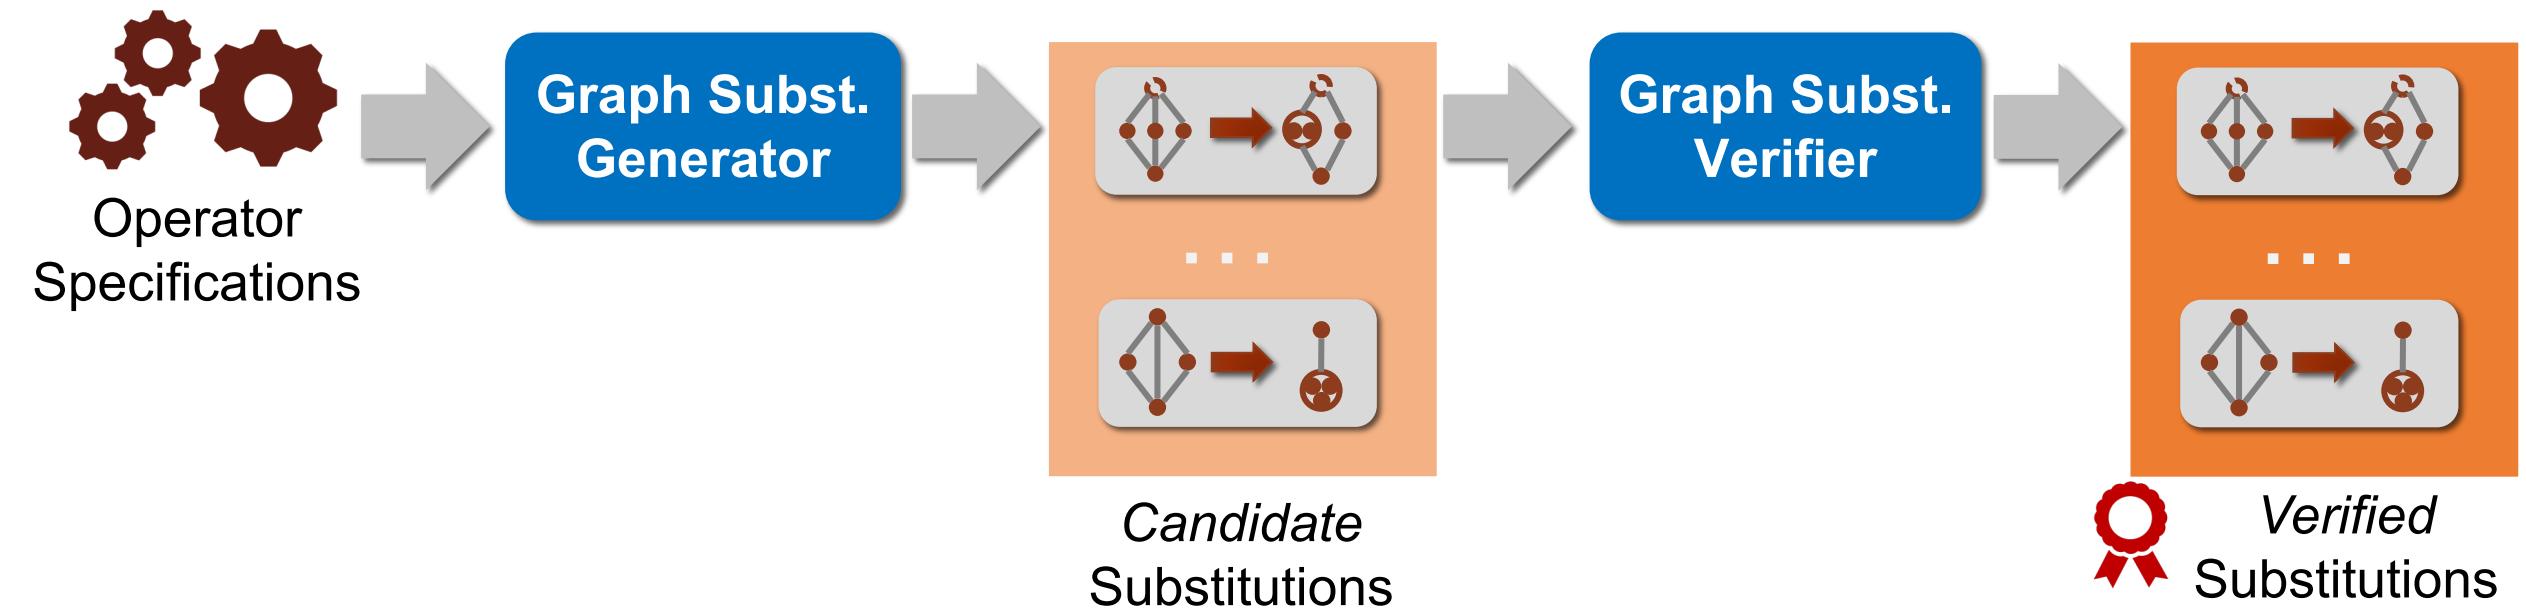
Only apply to **specific hardware**



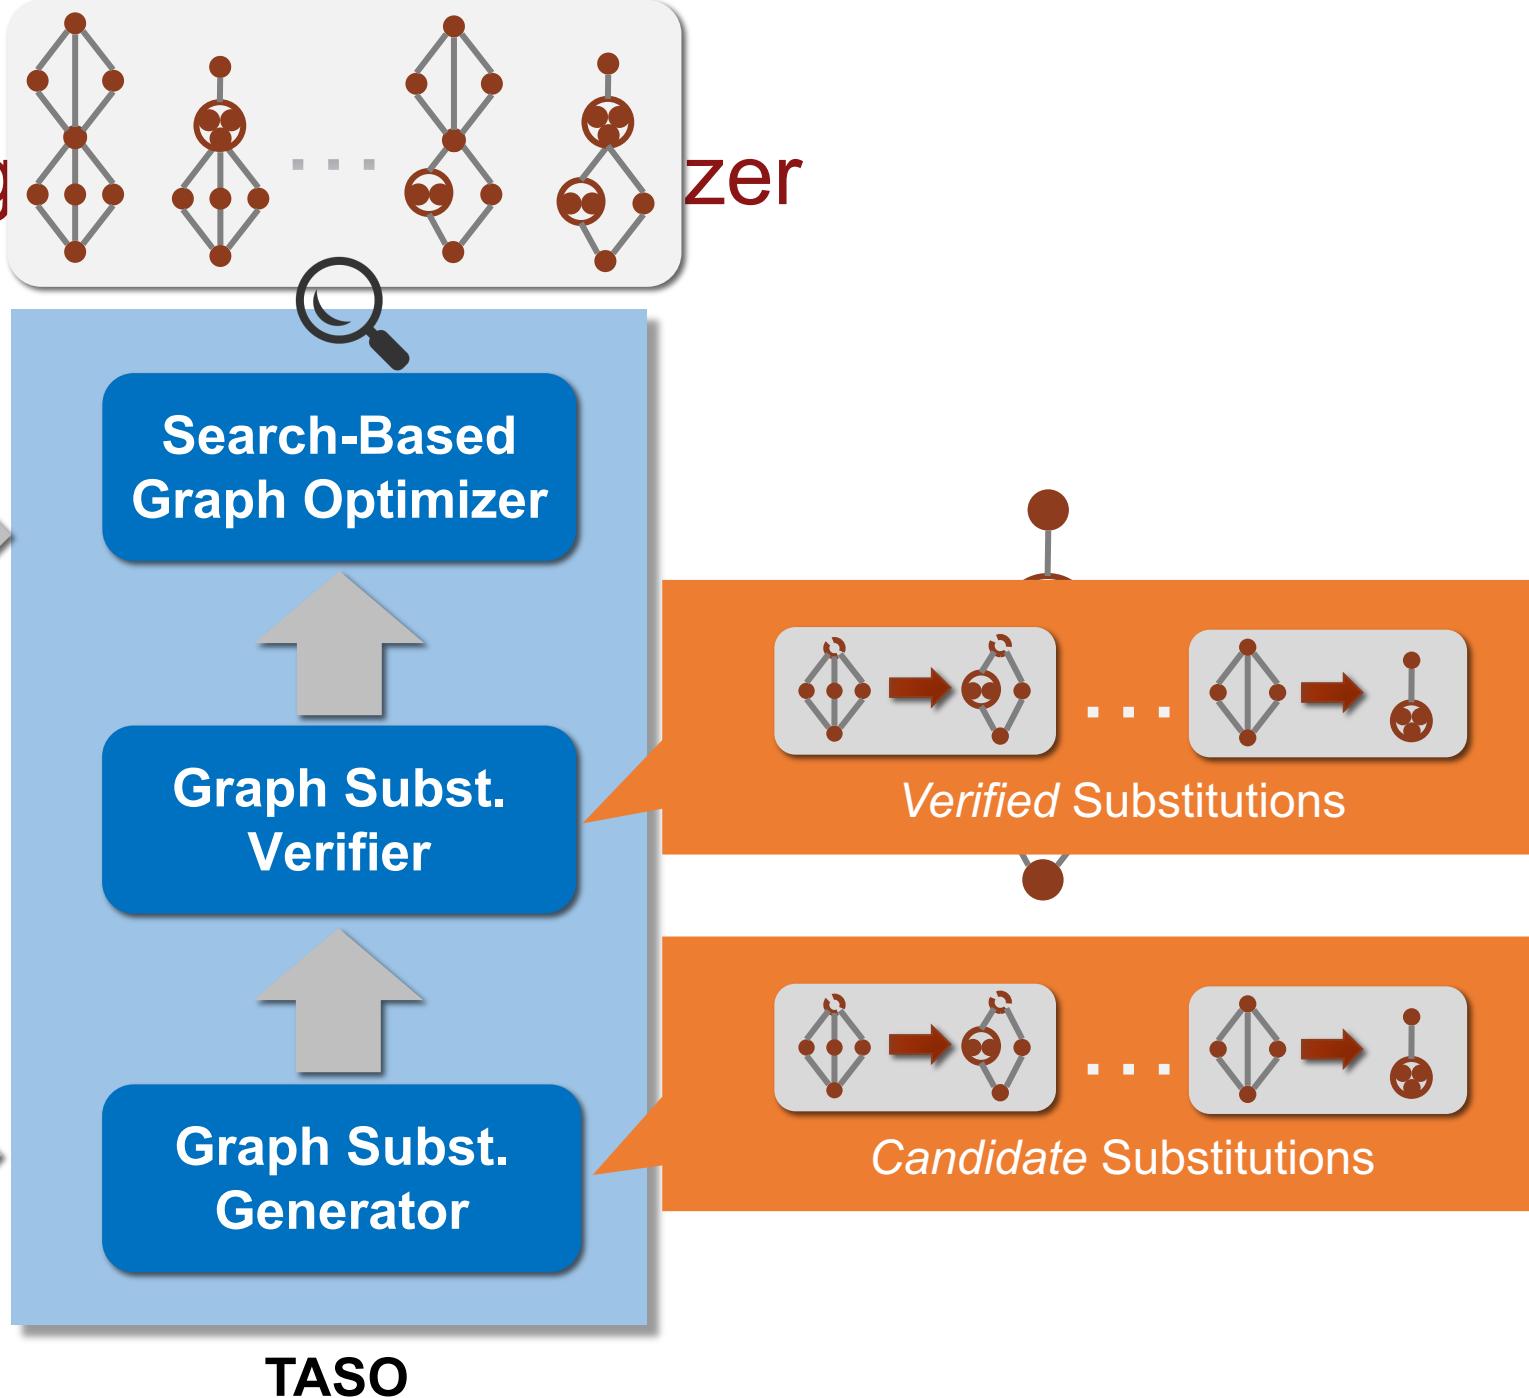
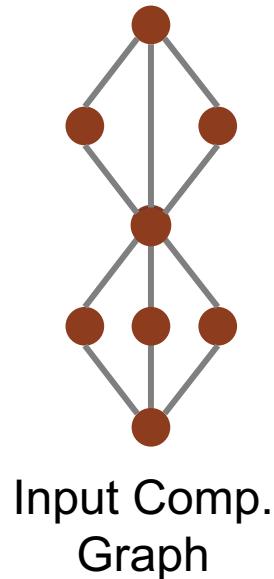
Only apply to **specialized graph structures**

TASO: Tensor Algebra SuperOptimizer

Key idea: automatically *generate* graph substitutions and *verify* them



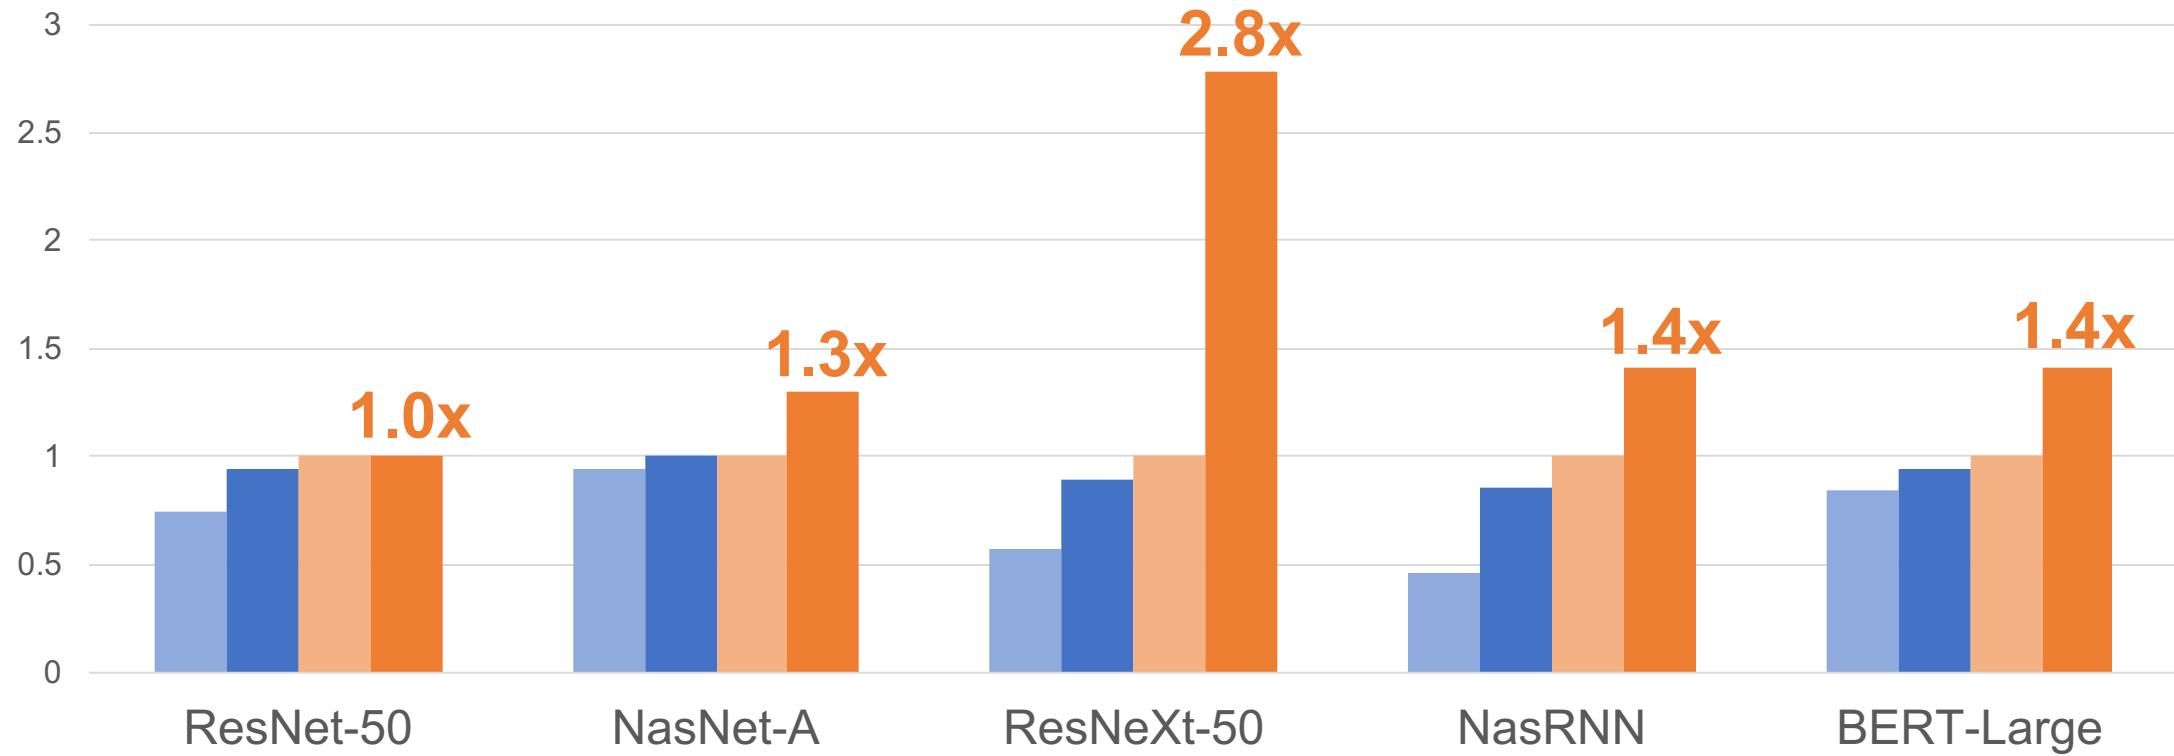
TASO: Tensor Alg Optimizer



End-to-end Inference Performance

Relative Speedup over Existing Frameworks

■ TensorFlow ■ TensorRT ■ MetaFlow ■ TASO



Joint Optimizer for Graph Substitution and Data Layout

- **Motivation:** some graph substitutions only improve performance when combined with particular layout transformations
- **Idea:** consider potential layout transformations along with graph substitutions
(additional 1.3x speedup)
- Cost-based backtracking search
 - Assume the cost to run a model is the sum of individual operators' costs
 - Measure the cost of each operator on hardware
 - A search takes less than **10 minutes**